Data Science with Python

Lesson 8—Machine Learning with Scikit–Learn

DATA SCIENCE

What You'll Learn

- What machine learning is and why it is important
- The machine learning approach
- Relevant terminologies that help you understand a dataset
- Features of supervised and unsupervised learning models
- Algorithms such as regression, classification, clustering, and dimensionality reduction

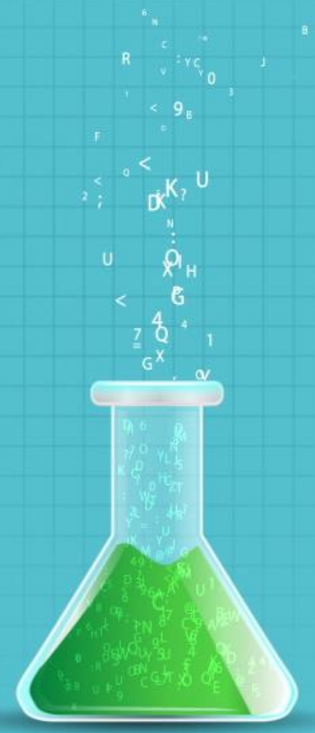

Why Machine Learning

If we stored the data generated in a day on Blu-ray disks and stacked them up, it would be equal to the height of four Eiffel towers! Machine learning helps analyze this data easily and quickly.

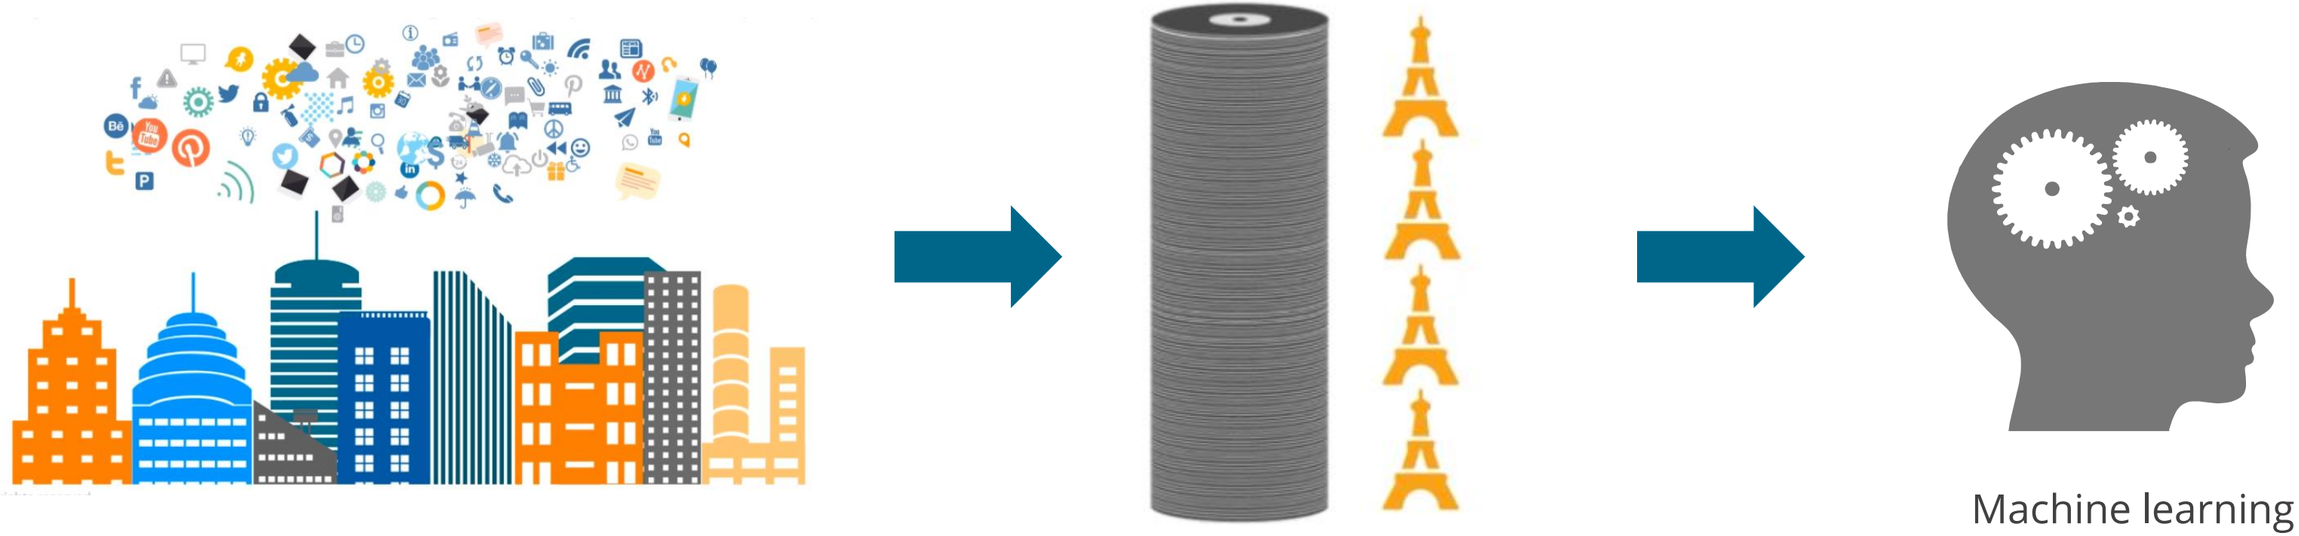

Machine learning

Purpose of Machine Learning

Machine learning is a great tool to analyze data, find hidden data patterns and relationships, and extract information to enable information-driven decisions and provide insights.

Data

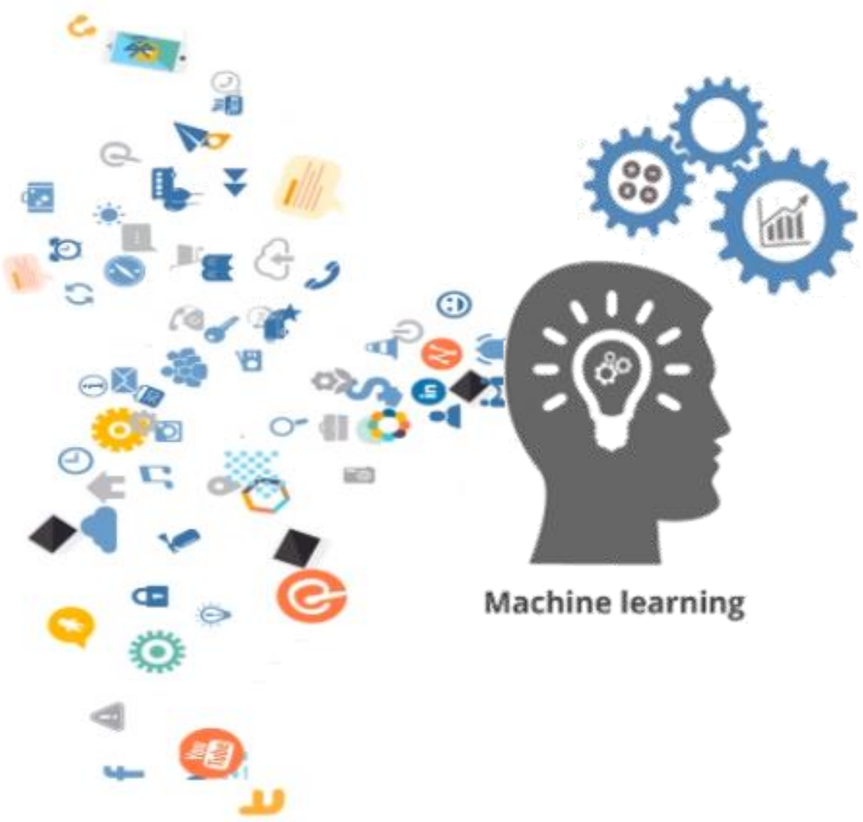

Identify patterns and relationships

Insights into unknown data

Information-driven decisions

Machine Learning Terminology

These are some machine learning terminologies that you will come across in this lesson:

Machine Learning Approach

The machine learning approach starts with either a problem that you need to solve or a given dataset that you need to analyze.

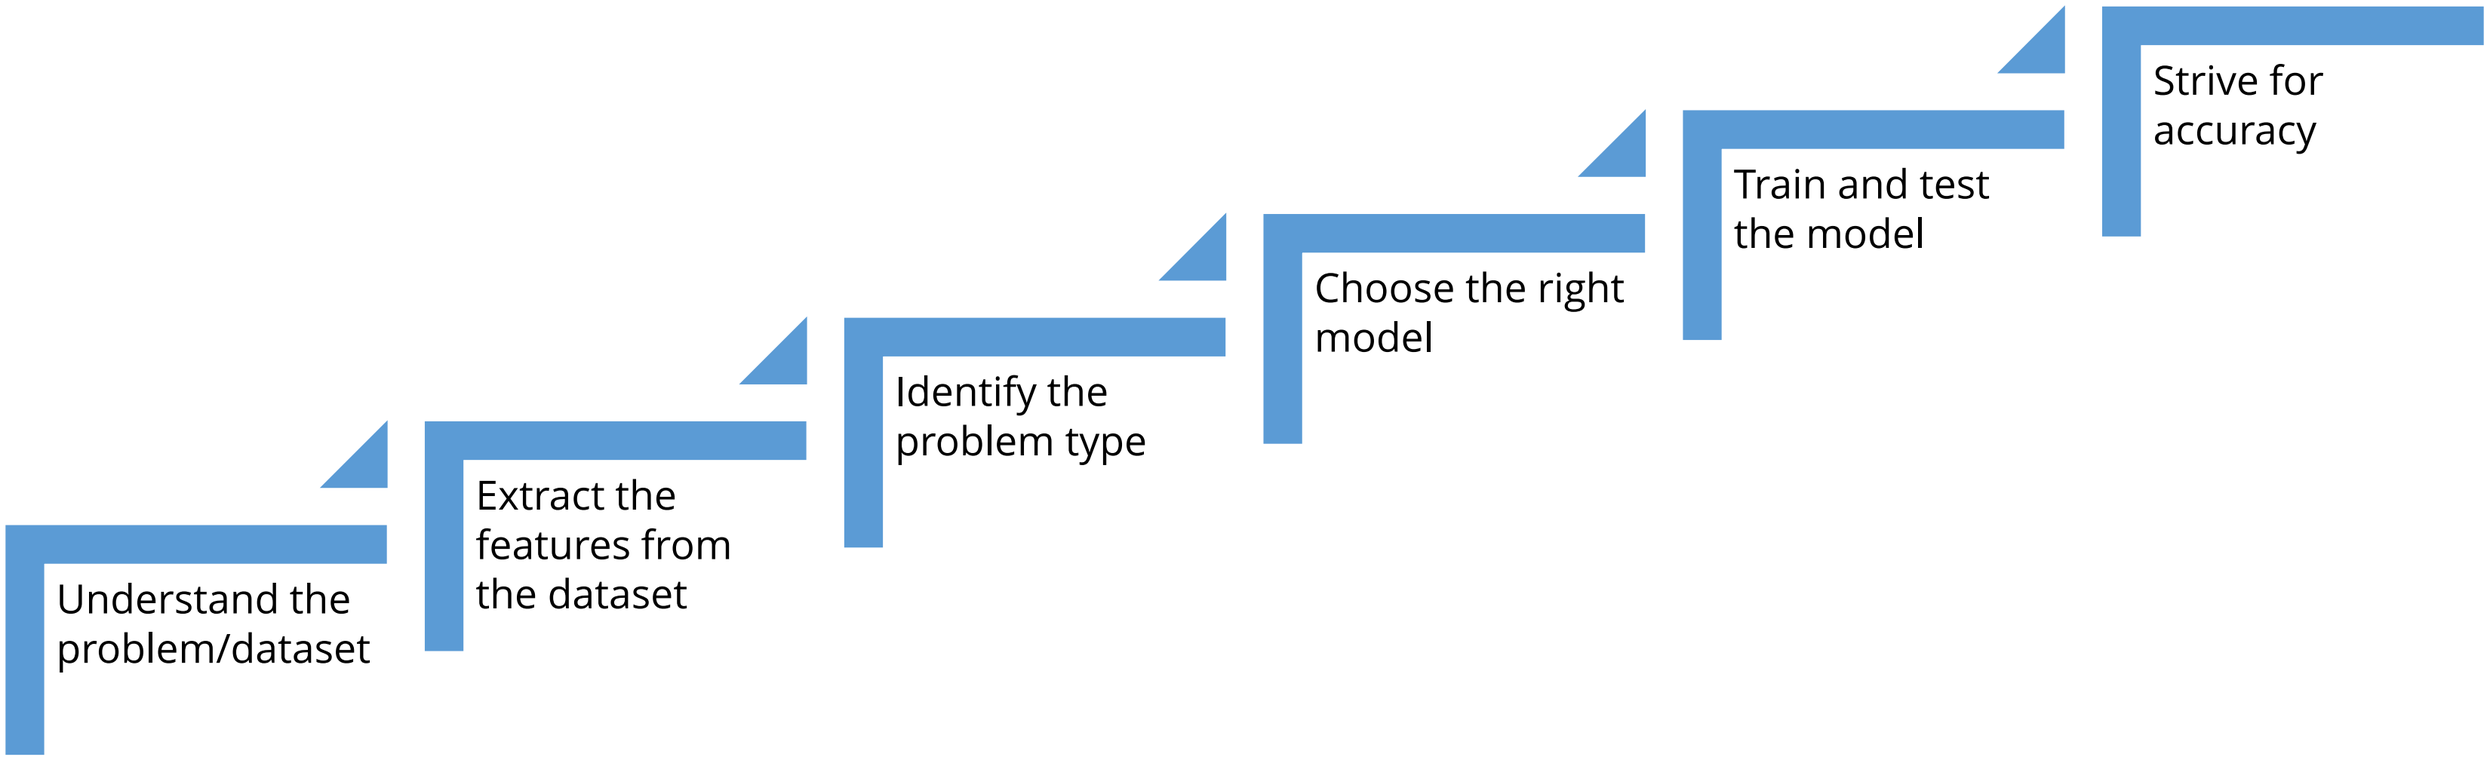

Steps 1 and 2: Understand the Dataset and Extract its Features

Let us look at a dataset and understand its features in terms of machine learning.

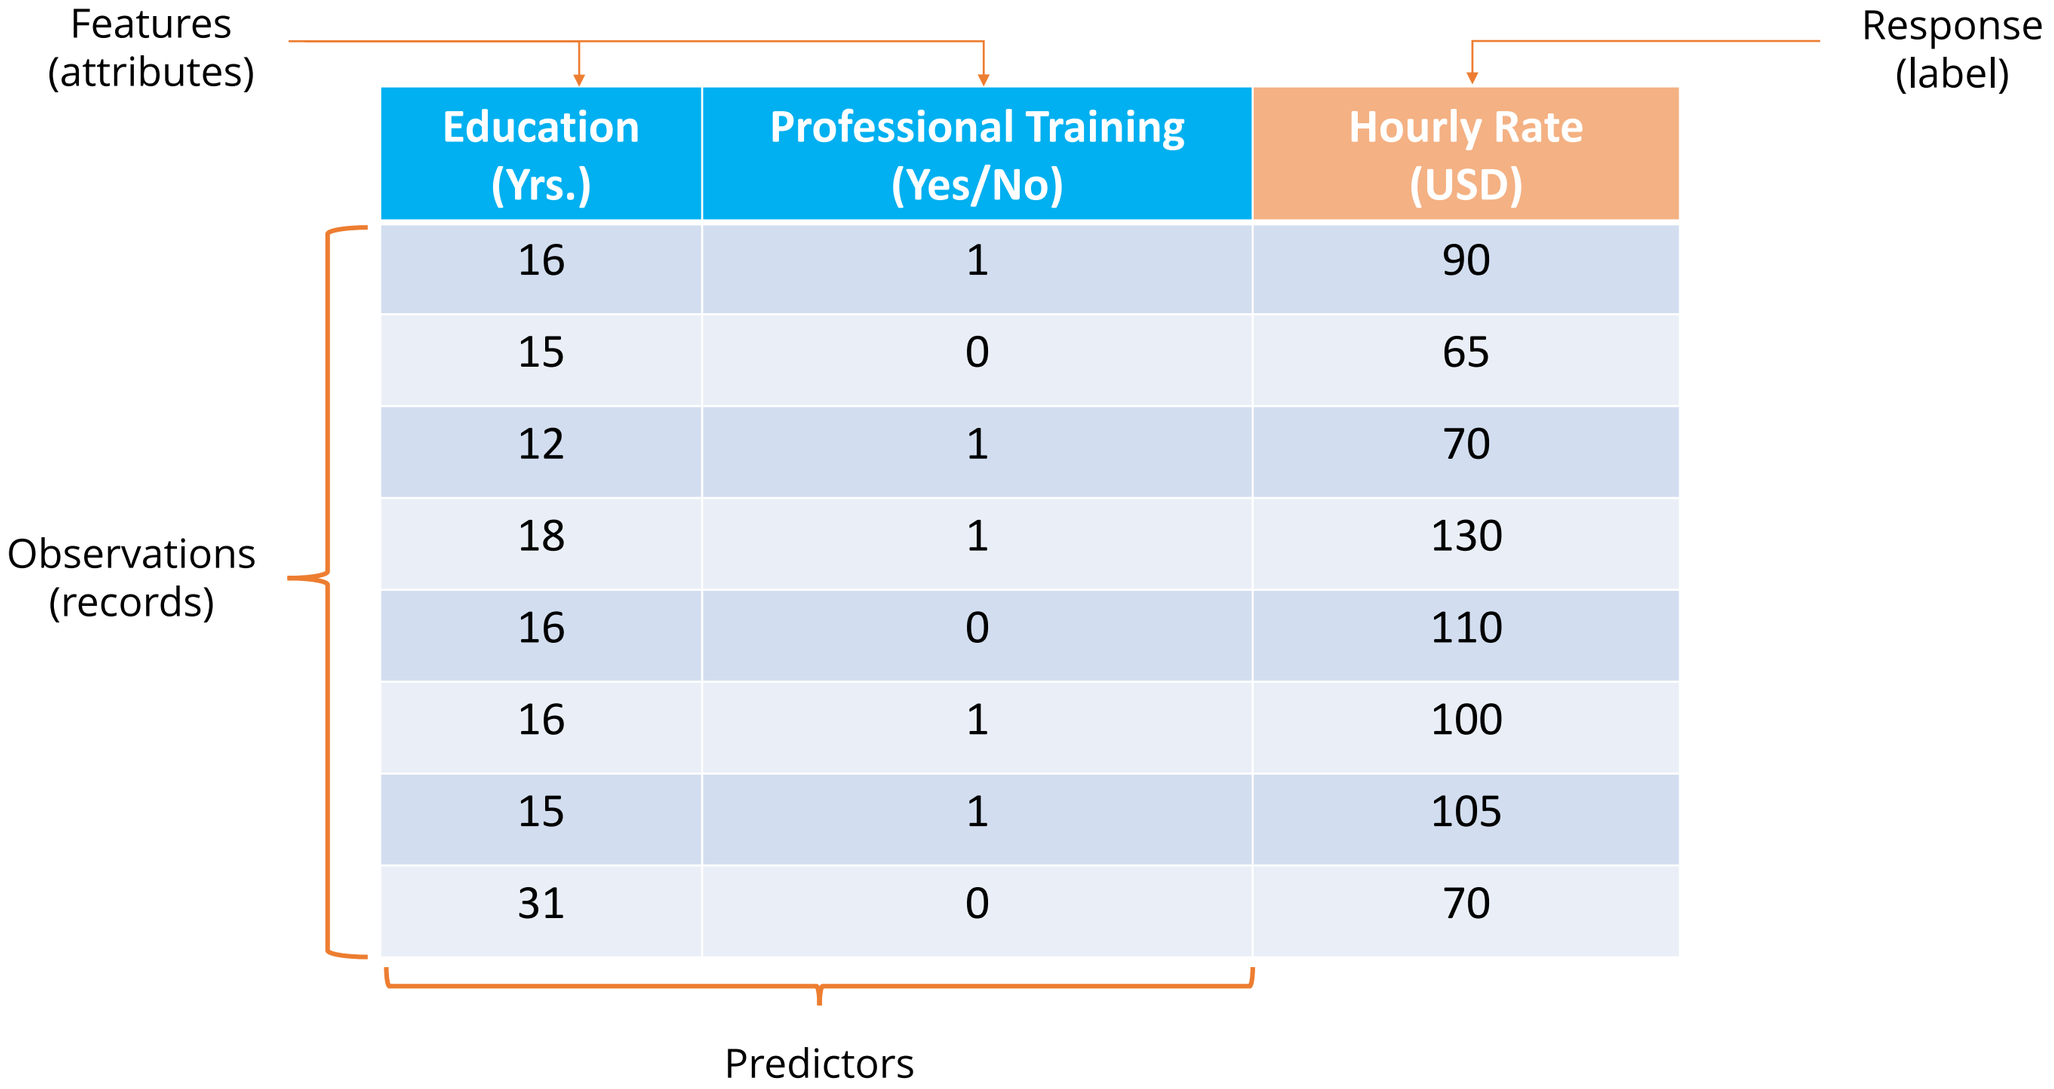

Education (Yrs.)	Professional Training (Yes/No)	Hourly Rate (USD)
16	1	90
15	0	65
12	1	70
18	1	130
16	0	110
16	1	100
15	1	105
31	0	70

Steps 3 and 4: Identify the Problem Type and Learning Model

Machine learning can either be supervised or unsupervised. The problem type should be selected based on the type of learning model.

Concept	Problem Types	Example

Supervised Learning

- In supervised learning, the dataset used to train a model should have observations, features, and responses. The model is trained to predict the “right” response for a given set of data points.
- Supervised learning models are used to predict an outcome.
- The goal of this model is to “generalize” a dataset so that the “general rule” can be applied to new data as well.

Unsupervised Learning

- In unsupervised learning, the response or the outcome of the data is not known.
- Supervised learning models are used to identify and visualize patterns in data by grouping similar types of data.
- The goal of this model is to “represent” data in a way that meaningful information can be extracted.

Steps 3 and 4: Identify the Problem Type and Learning Model (contd.)

Data can either be continuous or categorical. Based on whether it is supervised or unsupervised learning, the problem type will differ.

Concept | Problem Types | Example

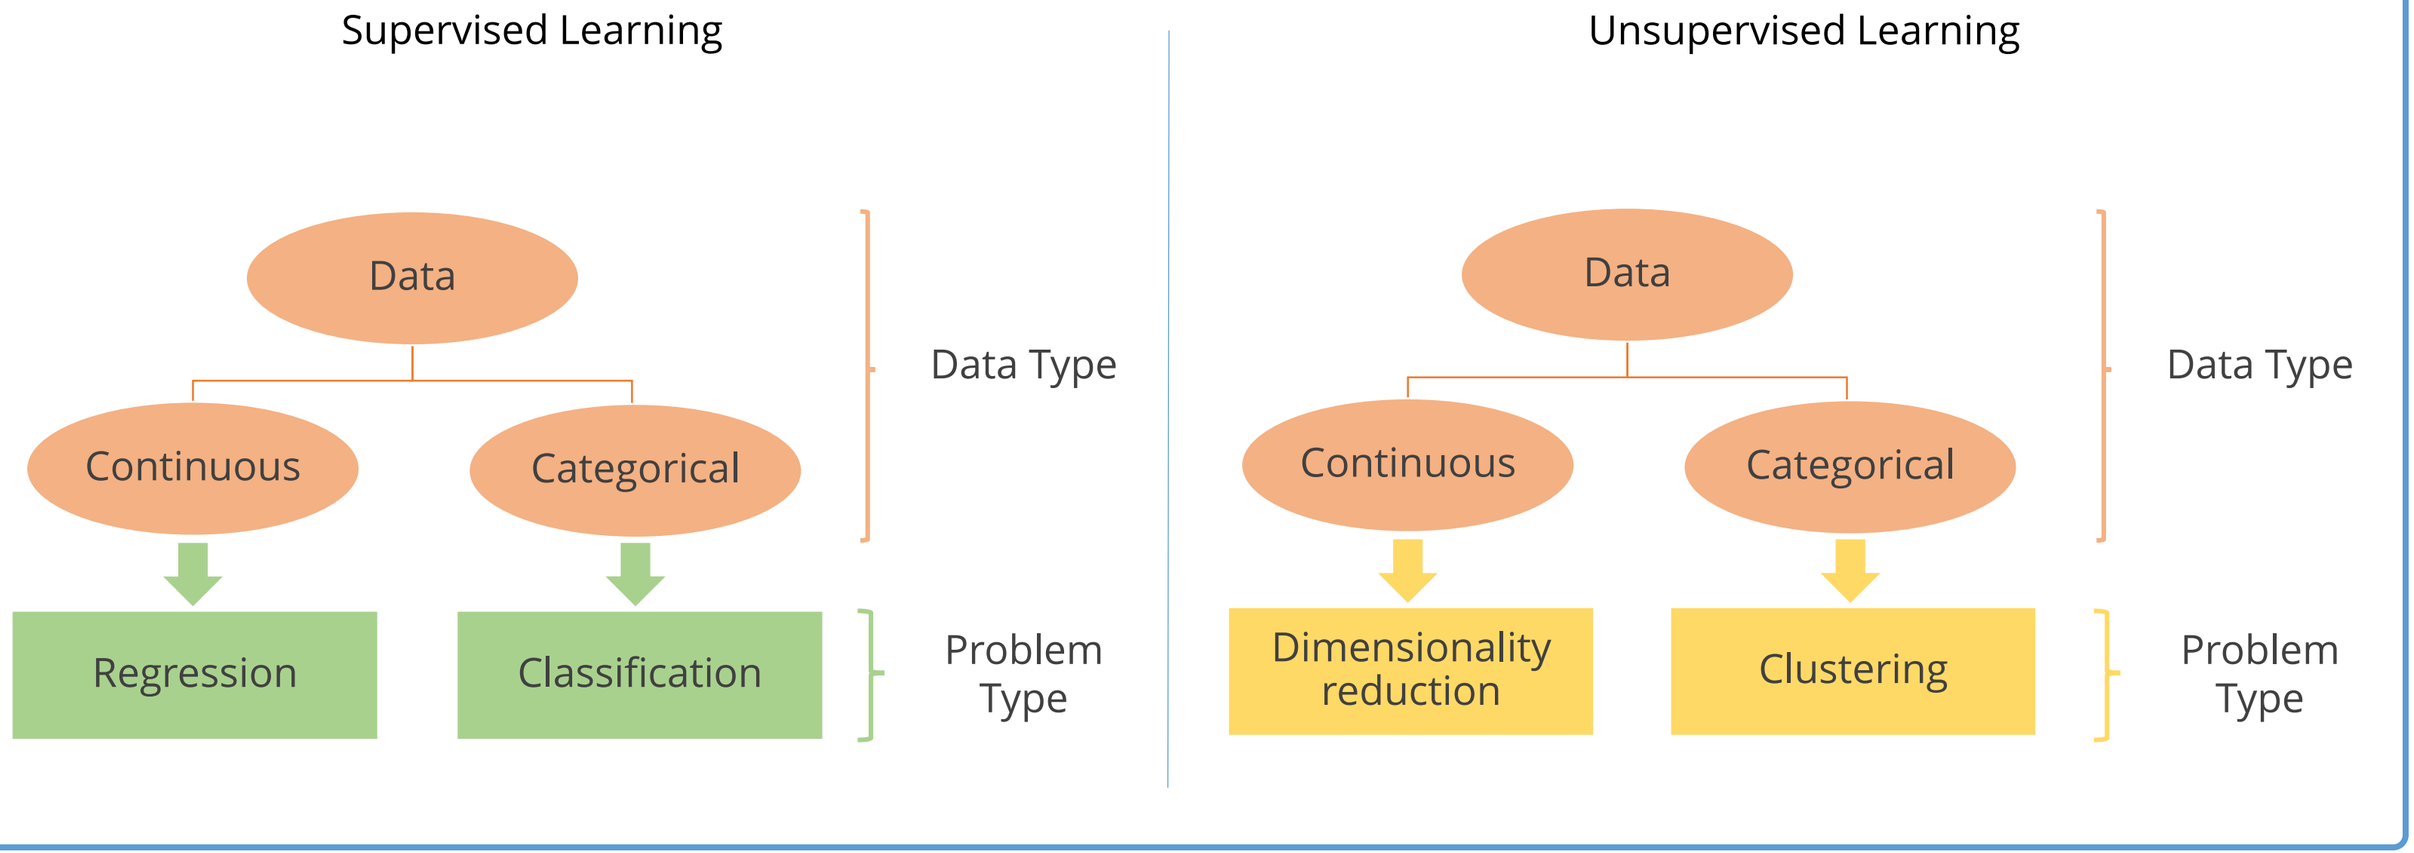

Steps 3 and 4: Identify the Problem Type and Learning Model (contd.)

Some examples of supervised and unsupervised learning models are shown here.

Concept | Problem Types | Example

Supervised Learning

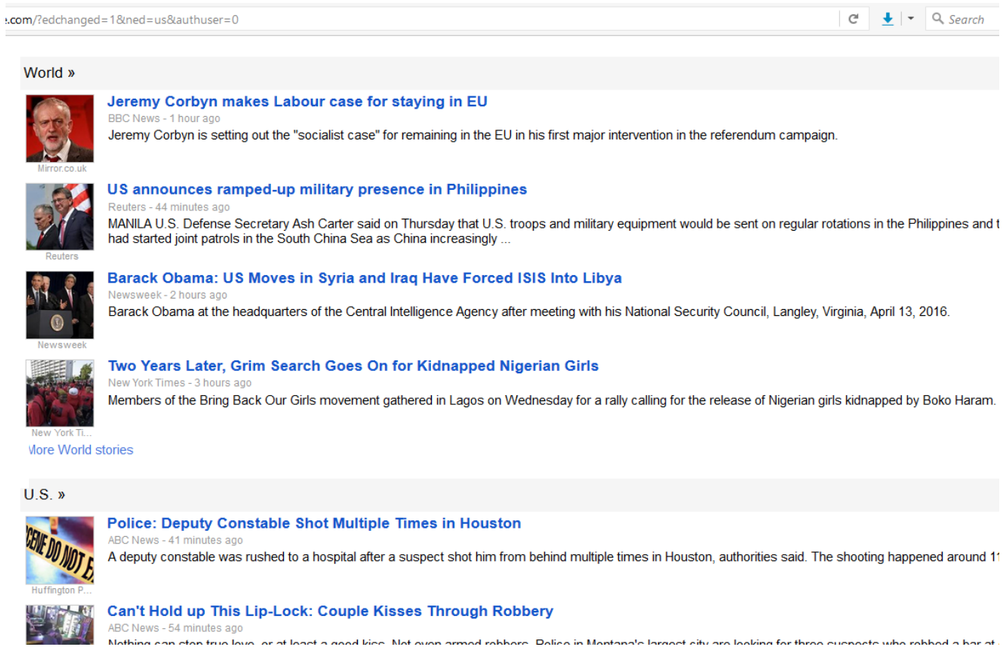

Categories of news based on the topics

Unsupervised Learning

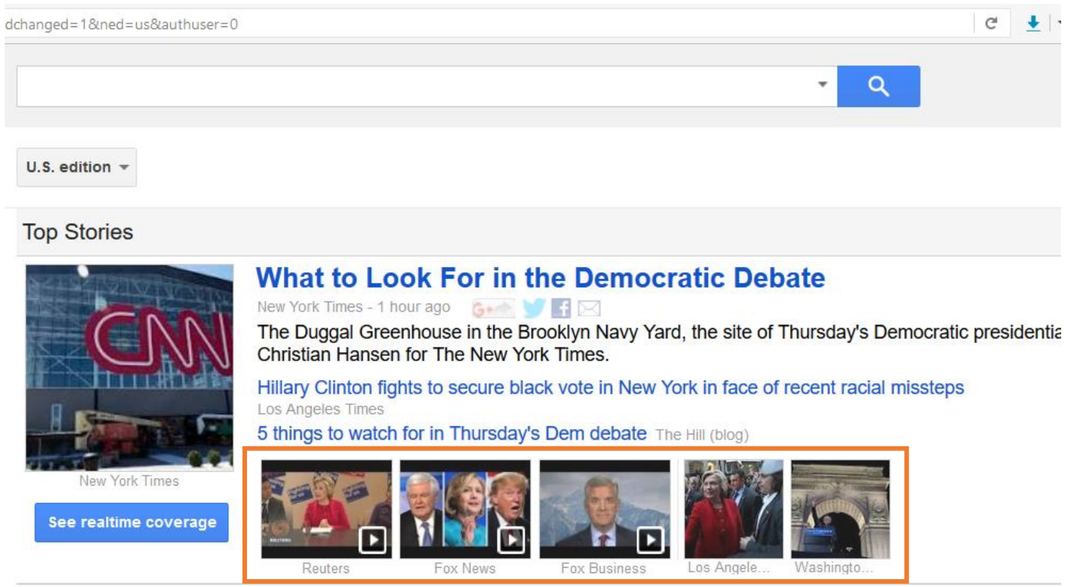

Grouping of similar stories on different news networks

How it Works—Supervised Learning Model

In supervised learning, a known dataset with observations, features, and response is used to create and train a machine learning algorithm. A predictive model, built on top of this algorithm, is then used to predict the response for a new dataset that has the same features.

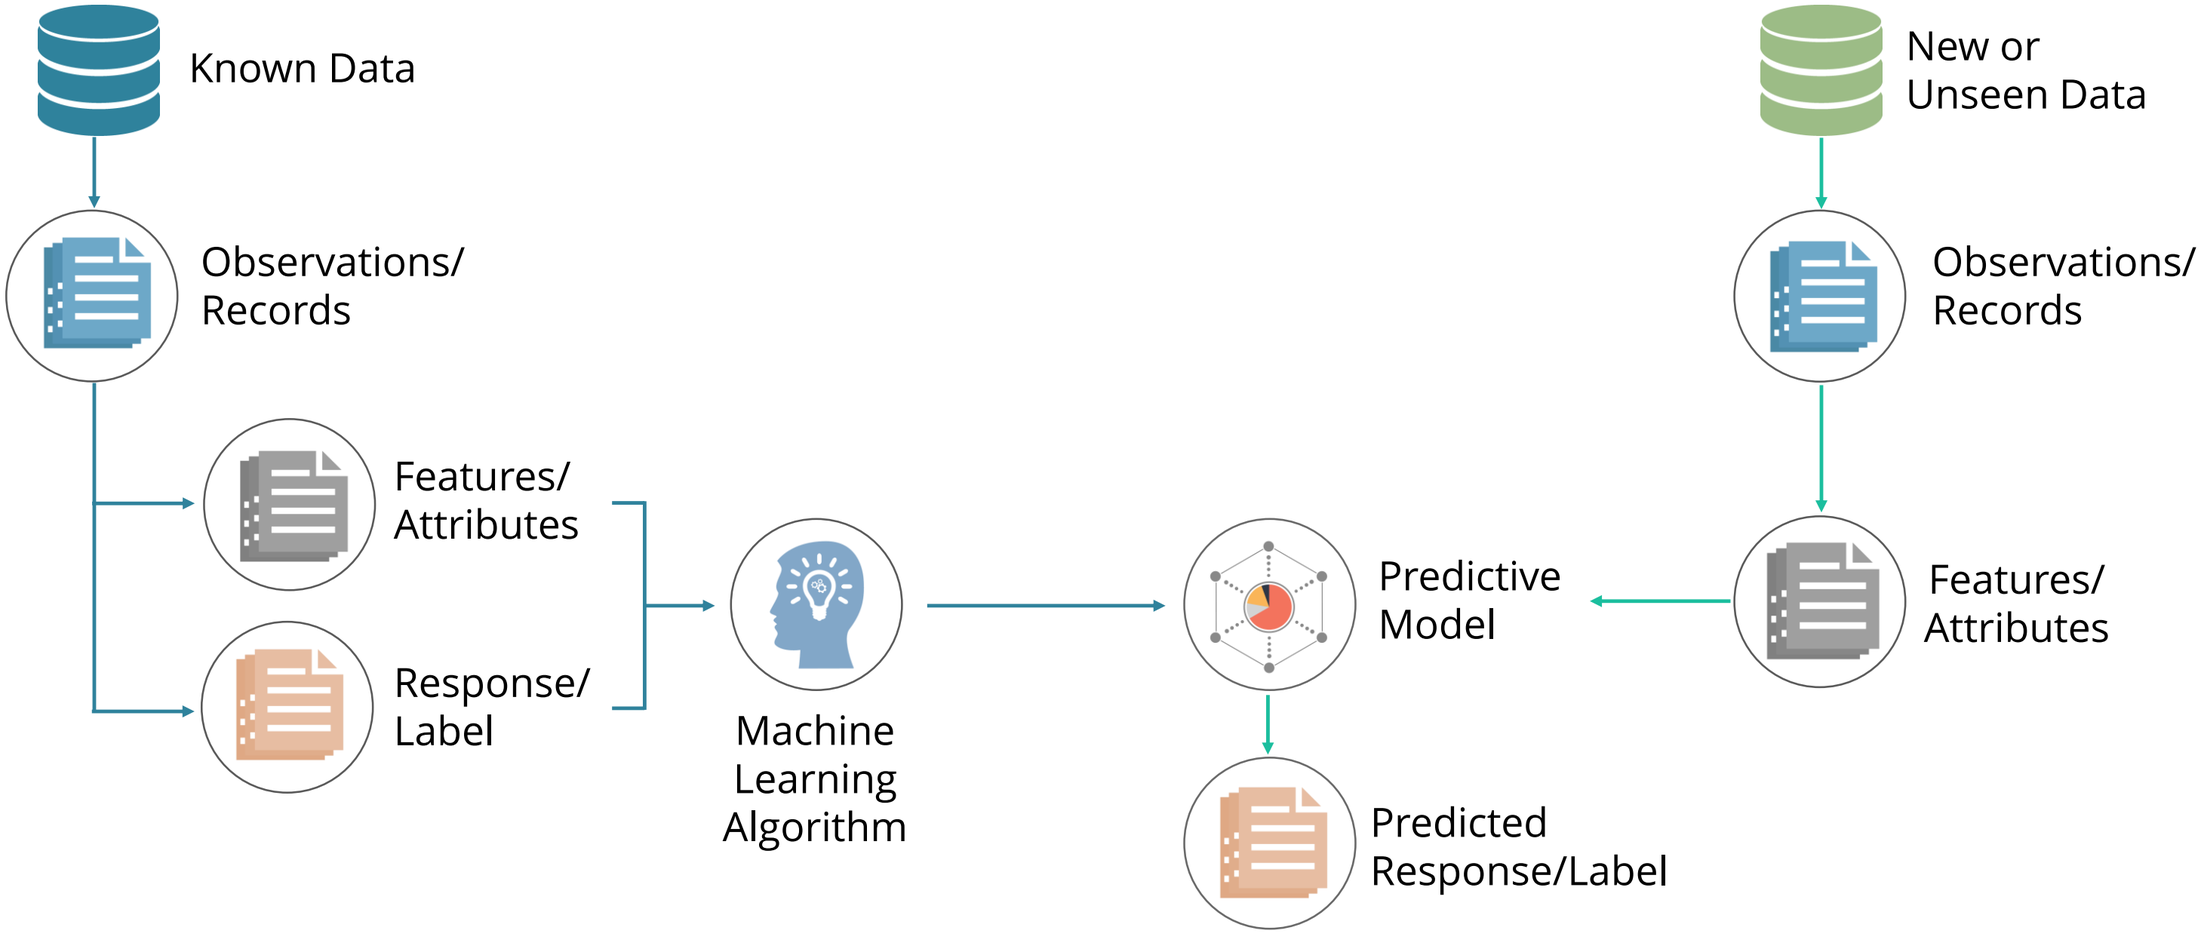

How it Works—Unsupervised Learning Model

In unsupervised learning, a known dataset has a set of observations with features. But the response is not known. The predictive model uses these features to identify how to classify and represent the data points of new or unseen data.

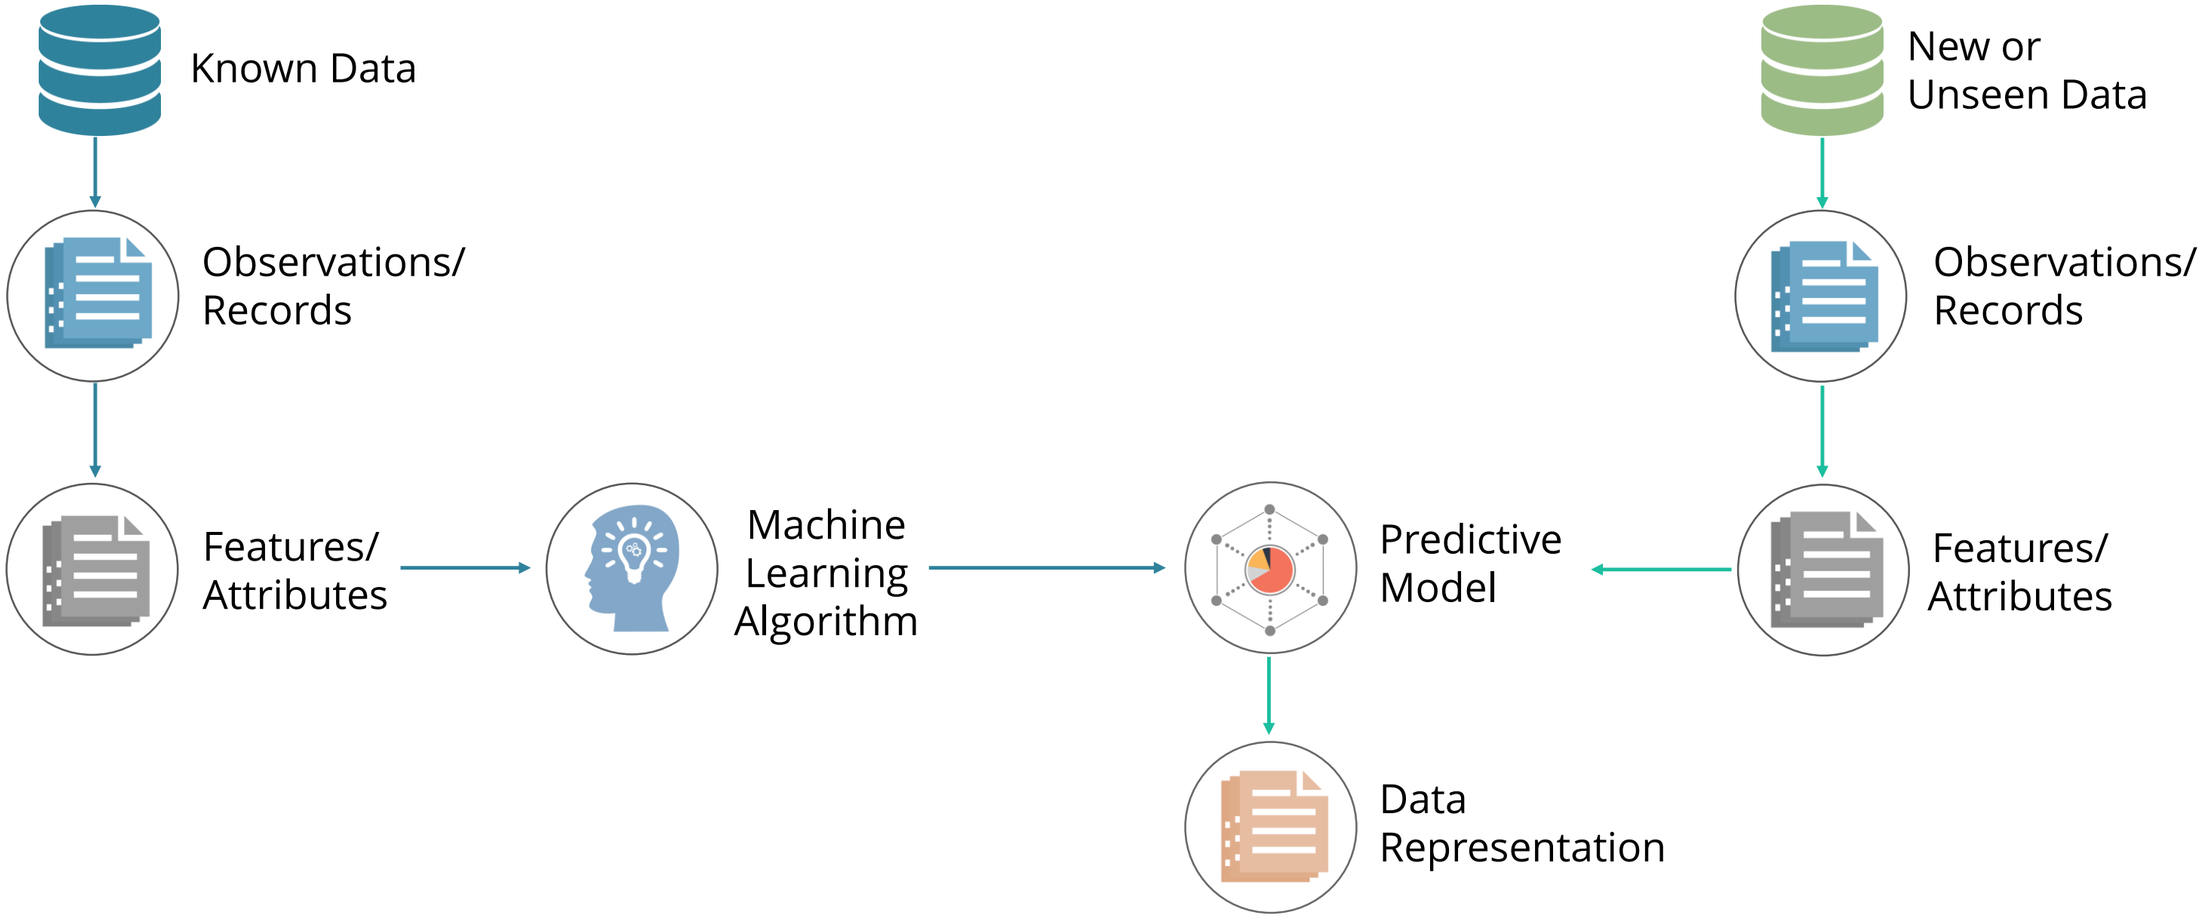

Steps 5 and 6: Train, Test, and Optimize the Model

To train supervised learning models, data analysts usually divide a known dataset into training and testing sets.

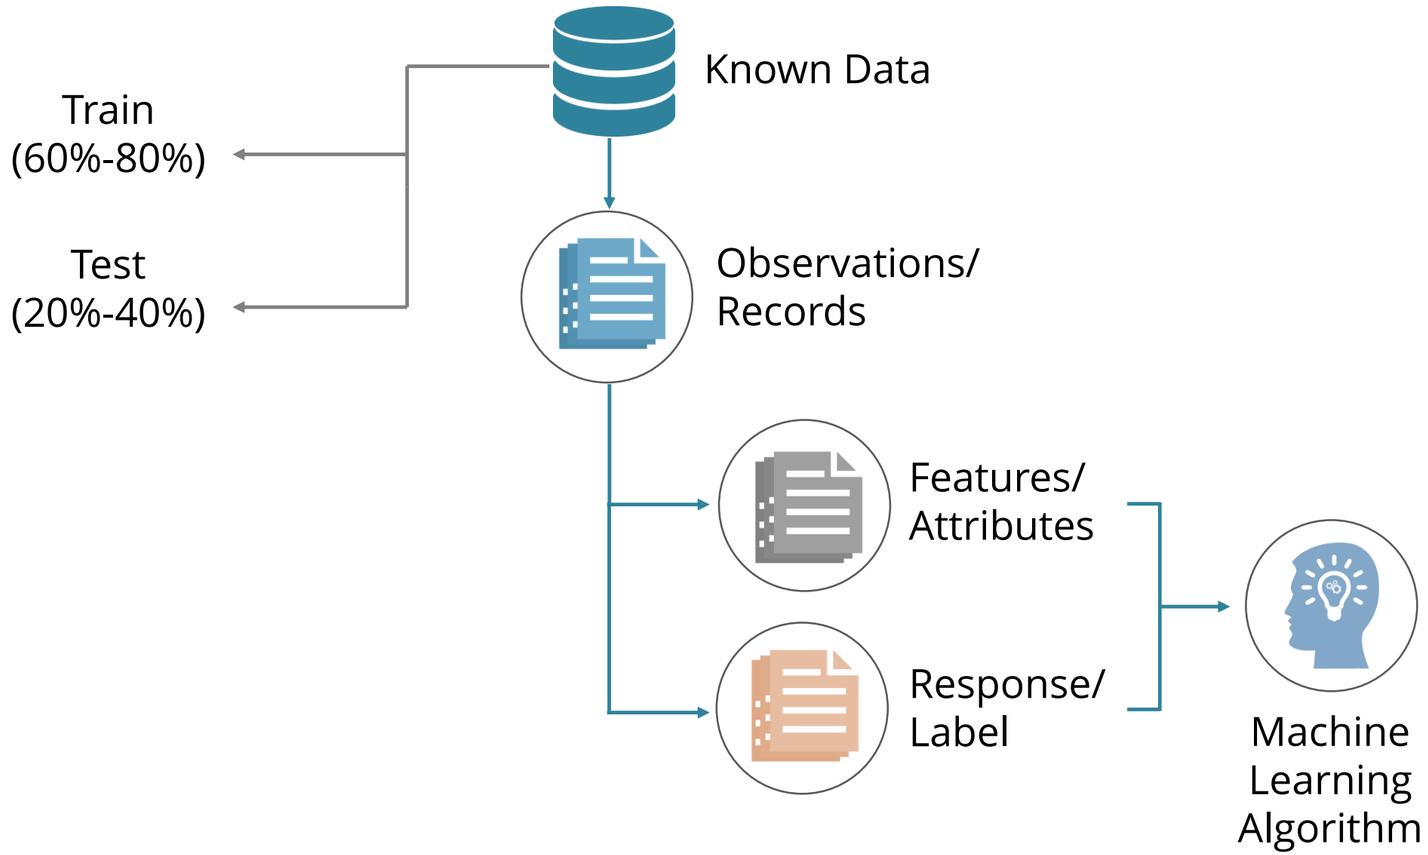

Steps 5 and 6: Train, Test, and Optimize the Model (contd.)

Let us look at an example to see how the split approach works.

Observation

Train set

Test set

Model Training

ID	Education	Professional training	Hourly rate
10	16	1	90
45	15	0	65
83	12	1	70
45	18	1	130
54	16	0	110
67	16	1	100
71	15	1	105
31	15	0	70

Response

Train set

Test set

Supervised Learning Model Considerations

Some considerations of supervised and unsupervised learning models are shown here.

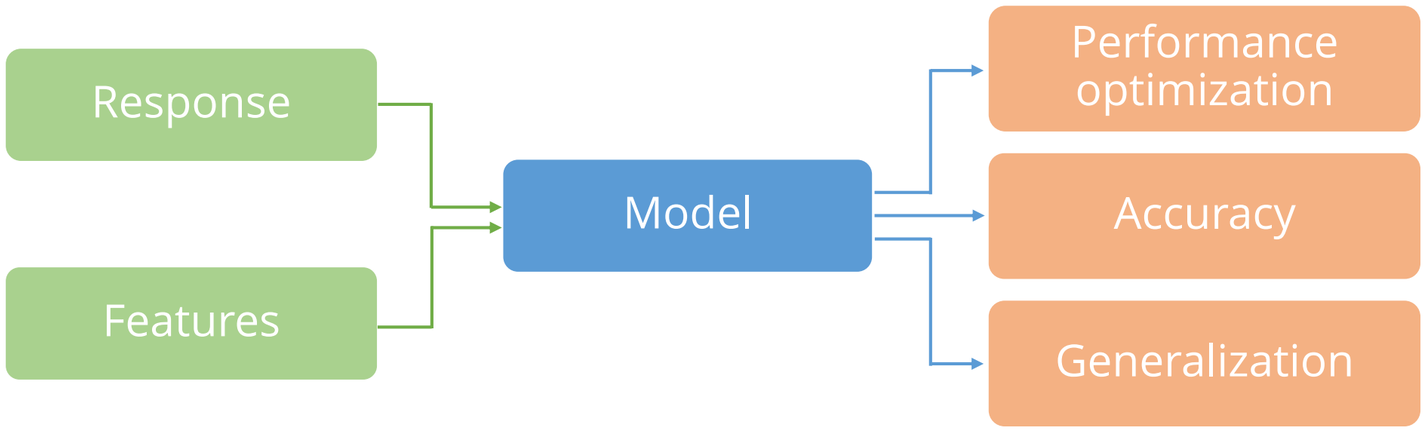

Knowledge Check

KNOWLEDGE CHECK

In machine learning, which one of the following is an observation?

a. Features

b. Attributes

c. Records

d. Labels

KNOWLEDGE CHECK

In machine learning, which one of the following is an observation?

a. Features

b. Attributes

c. Records

d. Labels

The correct answer is **c** .

Explanation: An observation is a set of examples, records, or samples.

KNOWLEDGE CHECK

If data is continuous and has labels (response), then it fits which of the following problem types?

a. Supervised learning: classification

b. Unsupervised learning: clustering

c. Unsupervised learning: dimensionality reduction

d. Supervised learning: regression

KNOWLEDGE CHECK

If data is continuous and has labels (response), then it fits which of the following problem types?

a. Supervised learning: classification

b. Unsupervised learning: clustering

c. Unsupervised learning: dimensionality reduction

d. Supervised learning: regression

The correct answer is **d**.

Explanation: The regression algorithm belonging to the supervised learning model is best suited to analyze continuous data.

KNOWLEDGE CHECK

Identify the goal of unsupervised learning. *Select all that apply*.

a. To predict the outcome

b. To understand the structure of the data

c. To generalize the dataset

d. To represent the data

KNOWLEDGE CHECK

Identify the goal of unsupervised learning. *Select all that apply*.

a. To predict the outcome

b. To understand the structure of the data

c. To generalize the dataset

d. To represent the data

The correct answer is **b, d**.

Explanation: The goal of unsupervised learning is to understand the structure of the data and represent it. There is no right or certain answer in unsupervised learning.

Scikit-Learn

Scikit is a powerful and modern machine learning Python library for fully and semi-automated data analysis and information extraction.

Efficient tools to identify and organize problems (Supervised/ Unsupervised)

Free and open datasets

Rich set of libraries for learning and predicting

Model support for every problem type

Model persistence

Open source community and vendor support

Scikit-Learn—Problem-Solution Approach

Scikit-learn helps Data Scientists organize their work through its problem-solution approach.

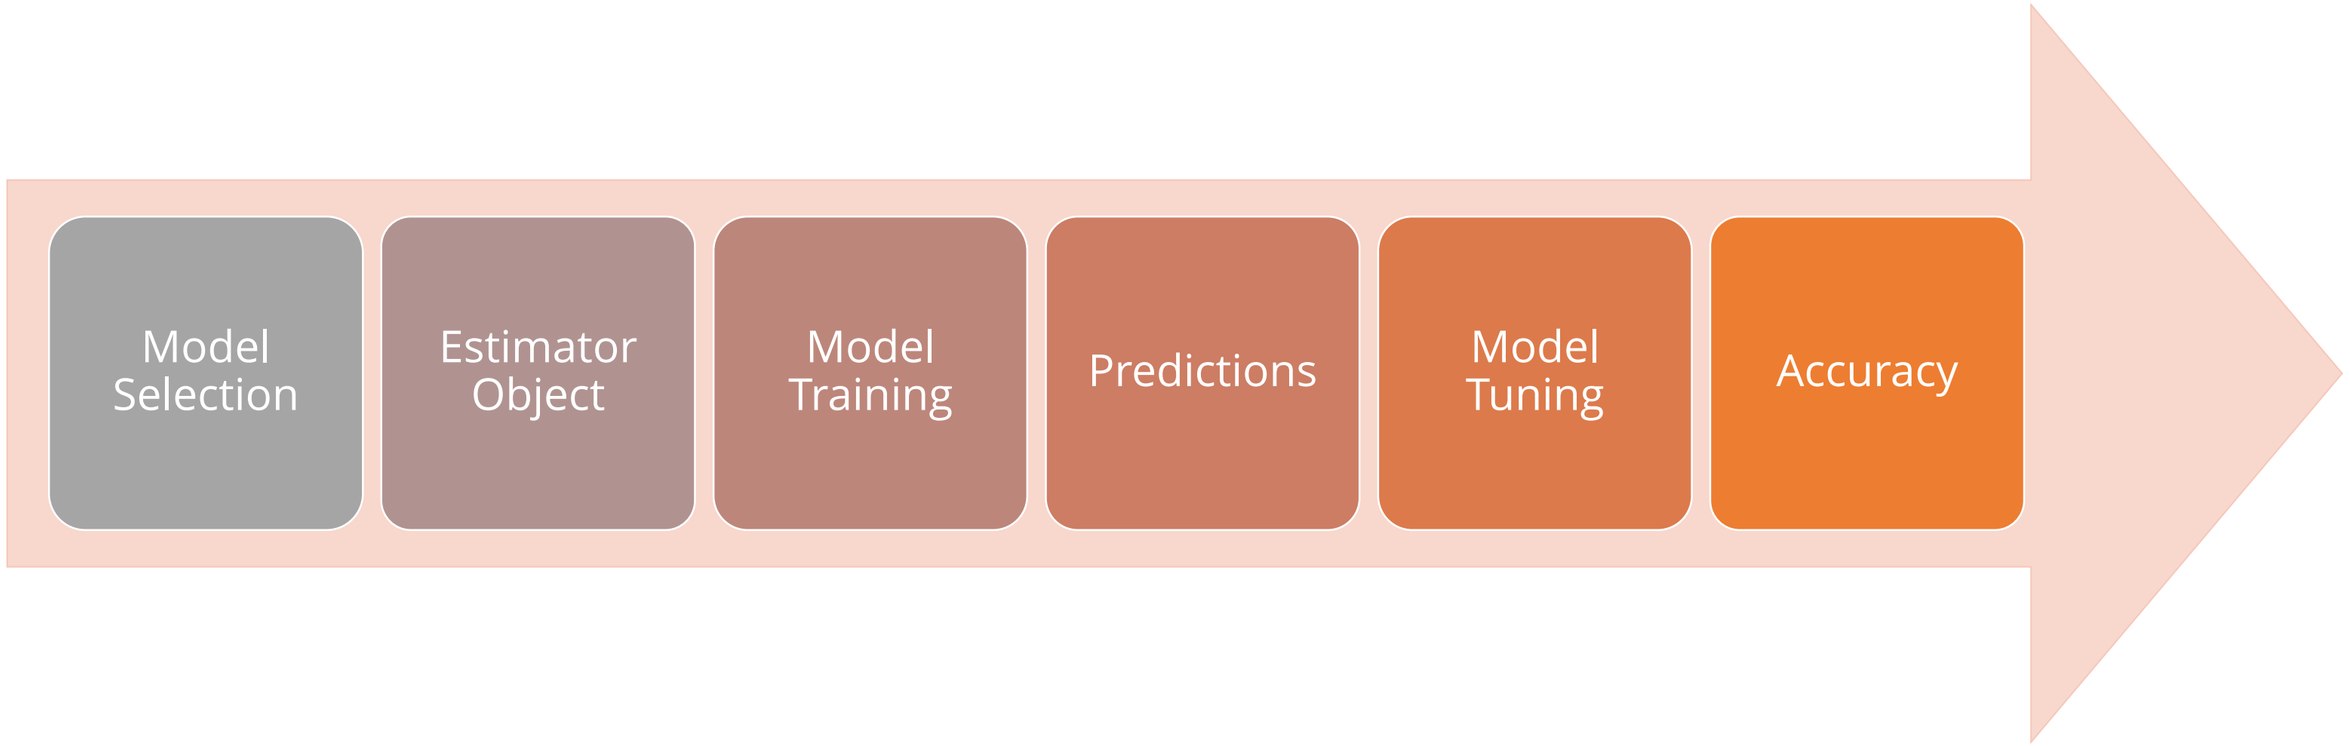

Scikit-Learn—Problem-Solution Considerations

While working with a Scikit-Learn dataset or loading your own data to Scikit -Learn, always consider these four points:

- Create separate objects for feature and response.
- Ensure that features and response have only numeric values.
- Features and response should be in the form of a NumPy ndarray.
- Since features and response would be in the form of arrays, they would have shapes and sizes.
- Features are always mapped as *x,* and response is mapped as *y*.

Knowledge Check

KNOWLEDGE CHECK

The estimator instance in Scikit-learn is a ____.

a. model

b. feature

c. dataset

d. response

KNOWLEDGE CHECK

The estimator instance in Scikit-learn is a _____.

a. model

b. feature

c. dataset

d. response

The correct answer is **a** .

Explanation: The estimator instance or object is a model.

Supervised Learning Models: Linear Regression

Linear regression is a supervised learning model used to analyze continuous data.

It is easy to use as the model does not require a lot of tuning.

It is the most basic and widely used technique to predict a value of an attribute.

It runs very fast, which makes it more time-efficient.

Supervised Learning Models: Linear Regression (contd.)

The linear regression equation is based on the formula for a simple linear equation.

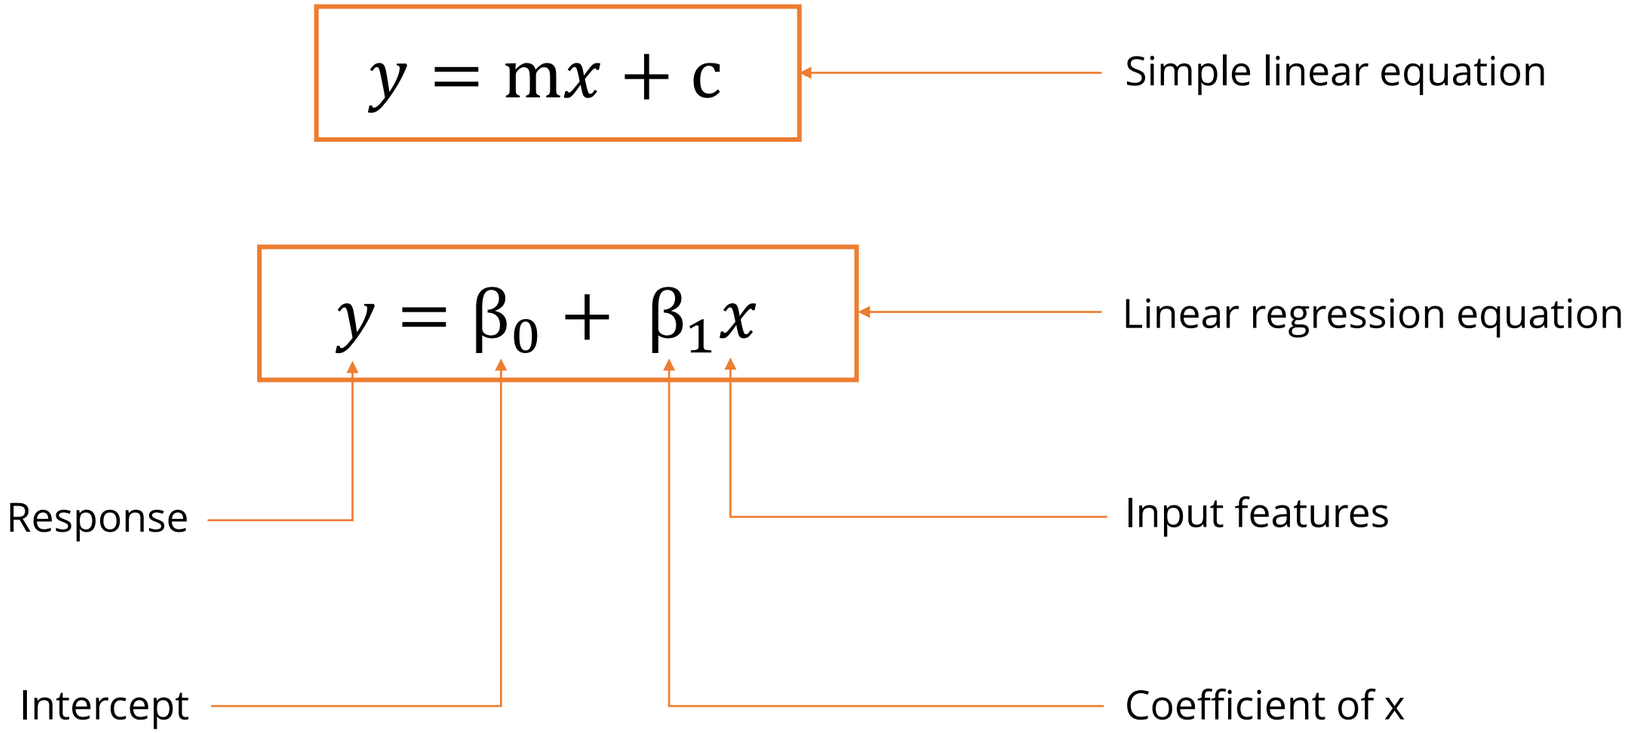

Supervised Learning Models: Linear Regression (contd.)

Linear regression is the most basic technique to predict a value of an attribute.

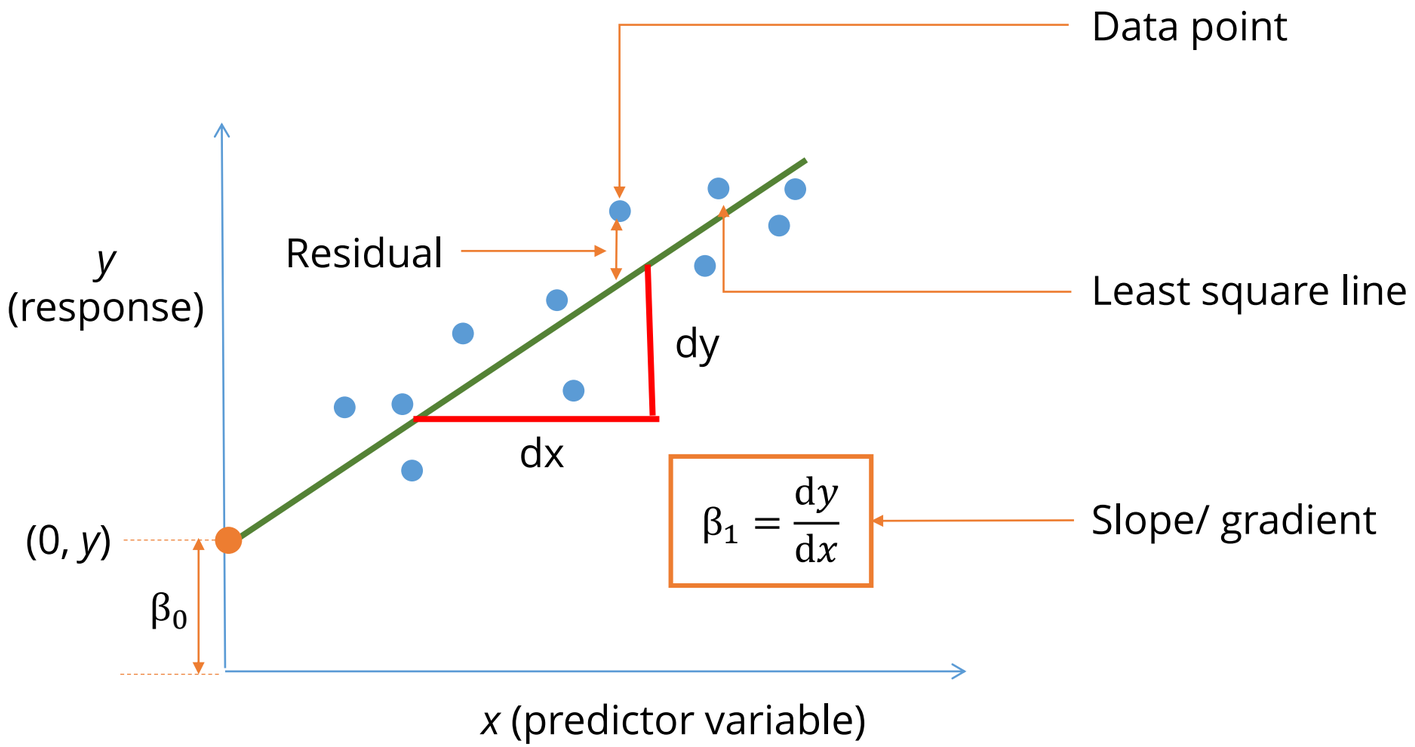

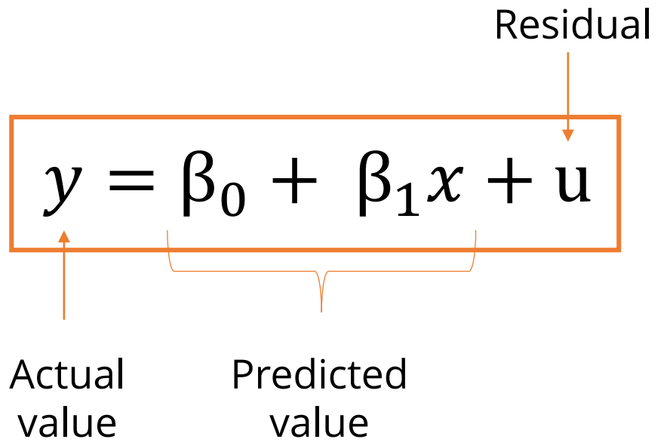

! The attributes are usually fitted using the “least square” approach.

Supervised Learning Models: Linear Regression (contd.)

Smaller the value of SSR or SSE, the more accurate the prediction will be, which would make the model the best fit.

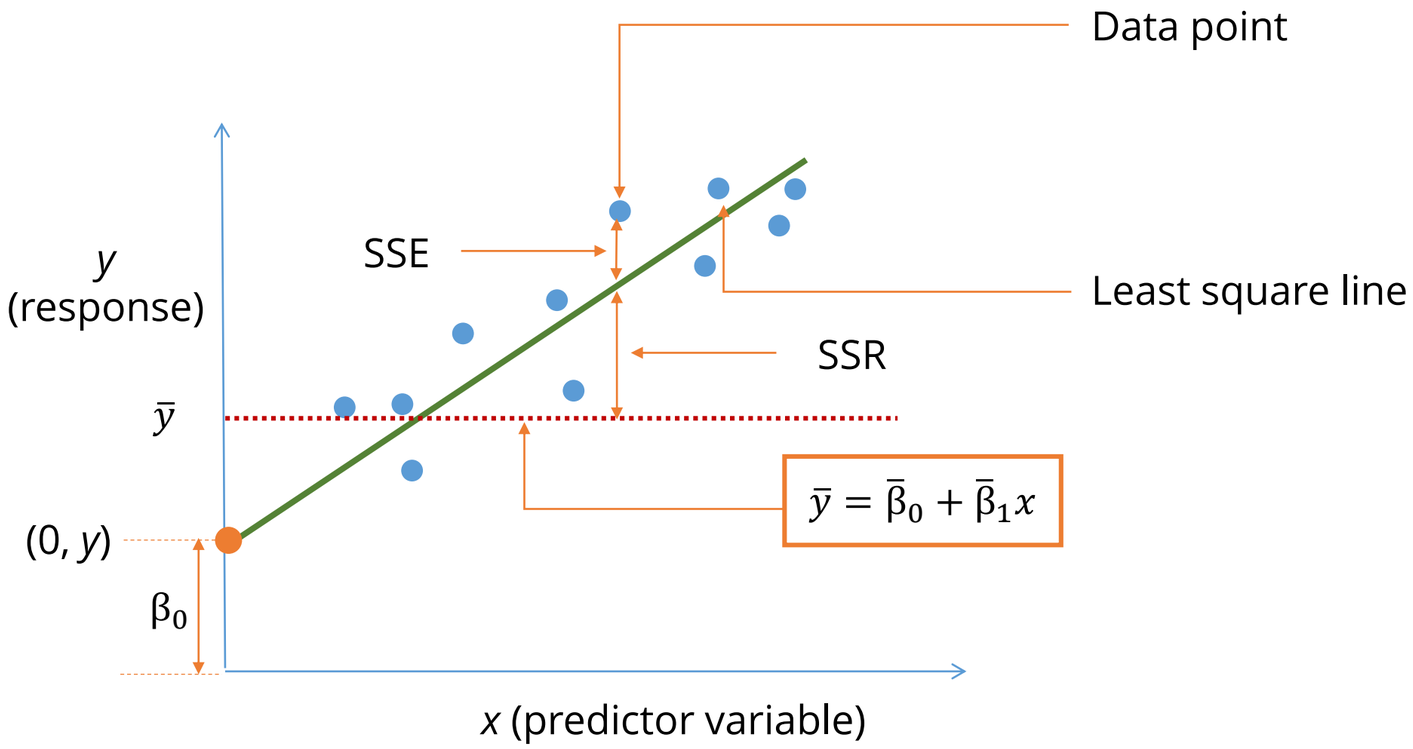

$$y = \beta_0 + \beta_1 x + u$$

$$SSR = \sum (\hat{y}_i - \bar{y})^2$$

Regression of sum of squares

$$SSE = \sum (y_i - \hat{y}_i)^2$$

Error of sum of squares

! The attributes are usually fitted using the “least square” approach.

Supervised Learning Models: Linear Regression (contd.)

Let us see how linear regression works in Scikit-Learn.

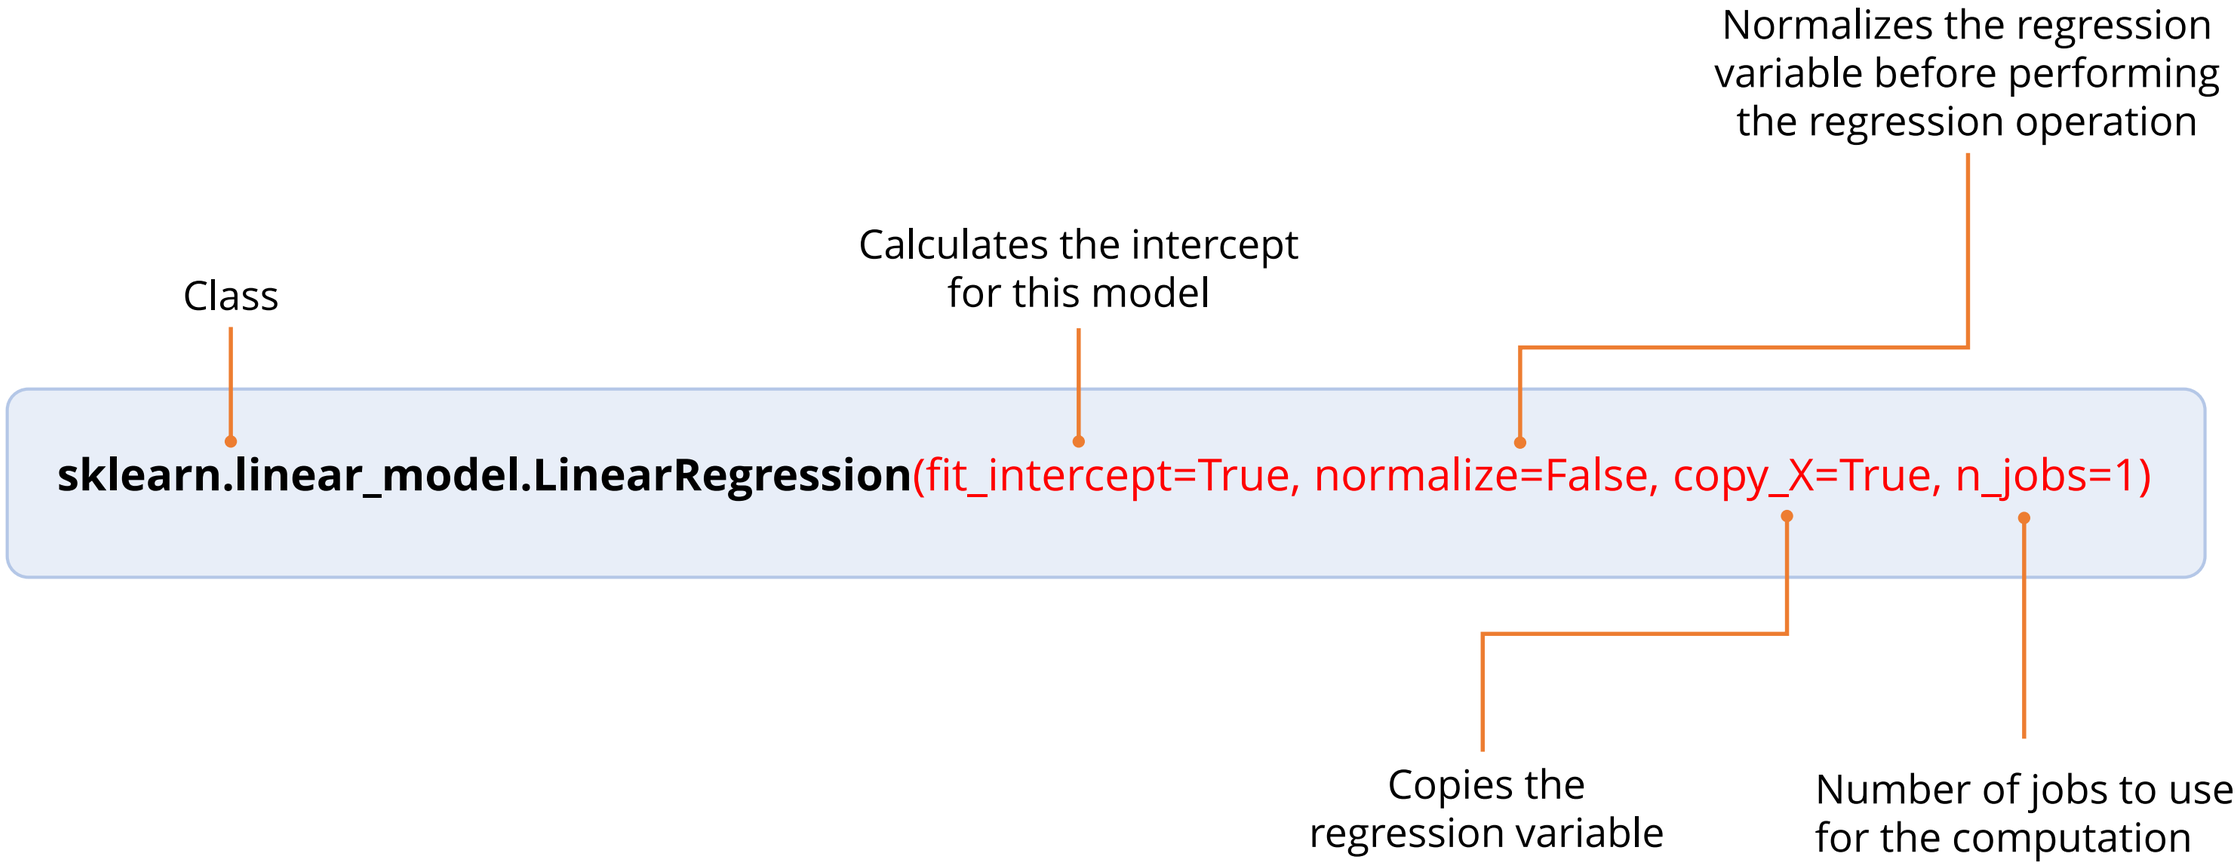

Demo 01—Loading a Dataset

Demonstrate how to load a built-in scikit-learn dataset

DATA SCIENCE

Demo 02—Linear Regression Model

Demonstrate how to create and train a linear regression model

DATA SCIENCE

Supervised Learning Models: Logistic Regression

Logistic regression is a generalization of the linear regression model used for classification problems.

$$\pi = \Pr(y = 1|x) = \frac{e^{\beta_0+\beta_1 x}}{1 + e^{\beta_0+\beta_1 x}}$$

Probability of *y* = 1, given *x*

Change in the log-odds for a unit change in x

The purpose of K-NN is to predict the class for each observation.

Supervised Learning Models: Logistic Regression (contd.)

Logistic regression is a generalization of the linear regression model used for classification problems.

$$\text{Odds} = \frac{\pi}{1-\pi}$$

Probability

$$\log\left(\frac{\pi}{1-\pi}\right) = \log\left(e^{\beta_0+\beta_1 x}\right) = \beta_0 + \beta_1 x$$

Logarithm of odds

Linear regression

Supervised Learning Models: Logistic Regression (contd.)

Logistic regression is a generalization of the linear regression model used for classification problems.

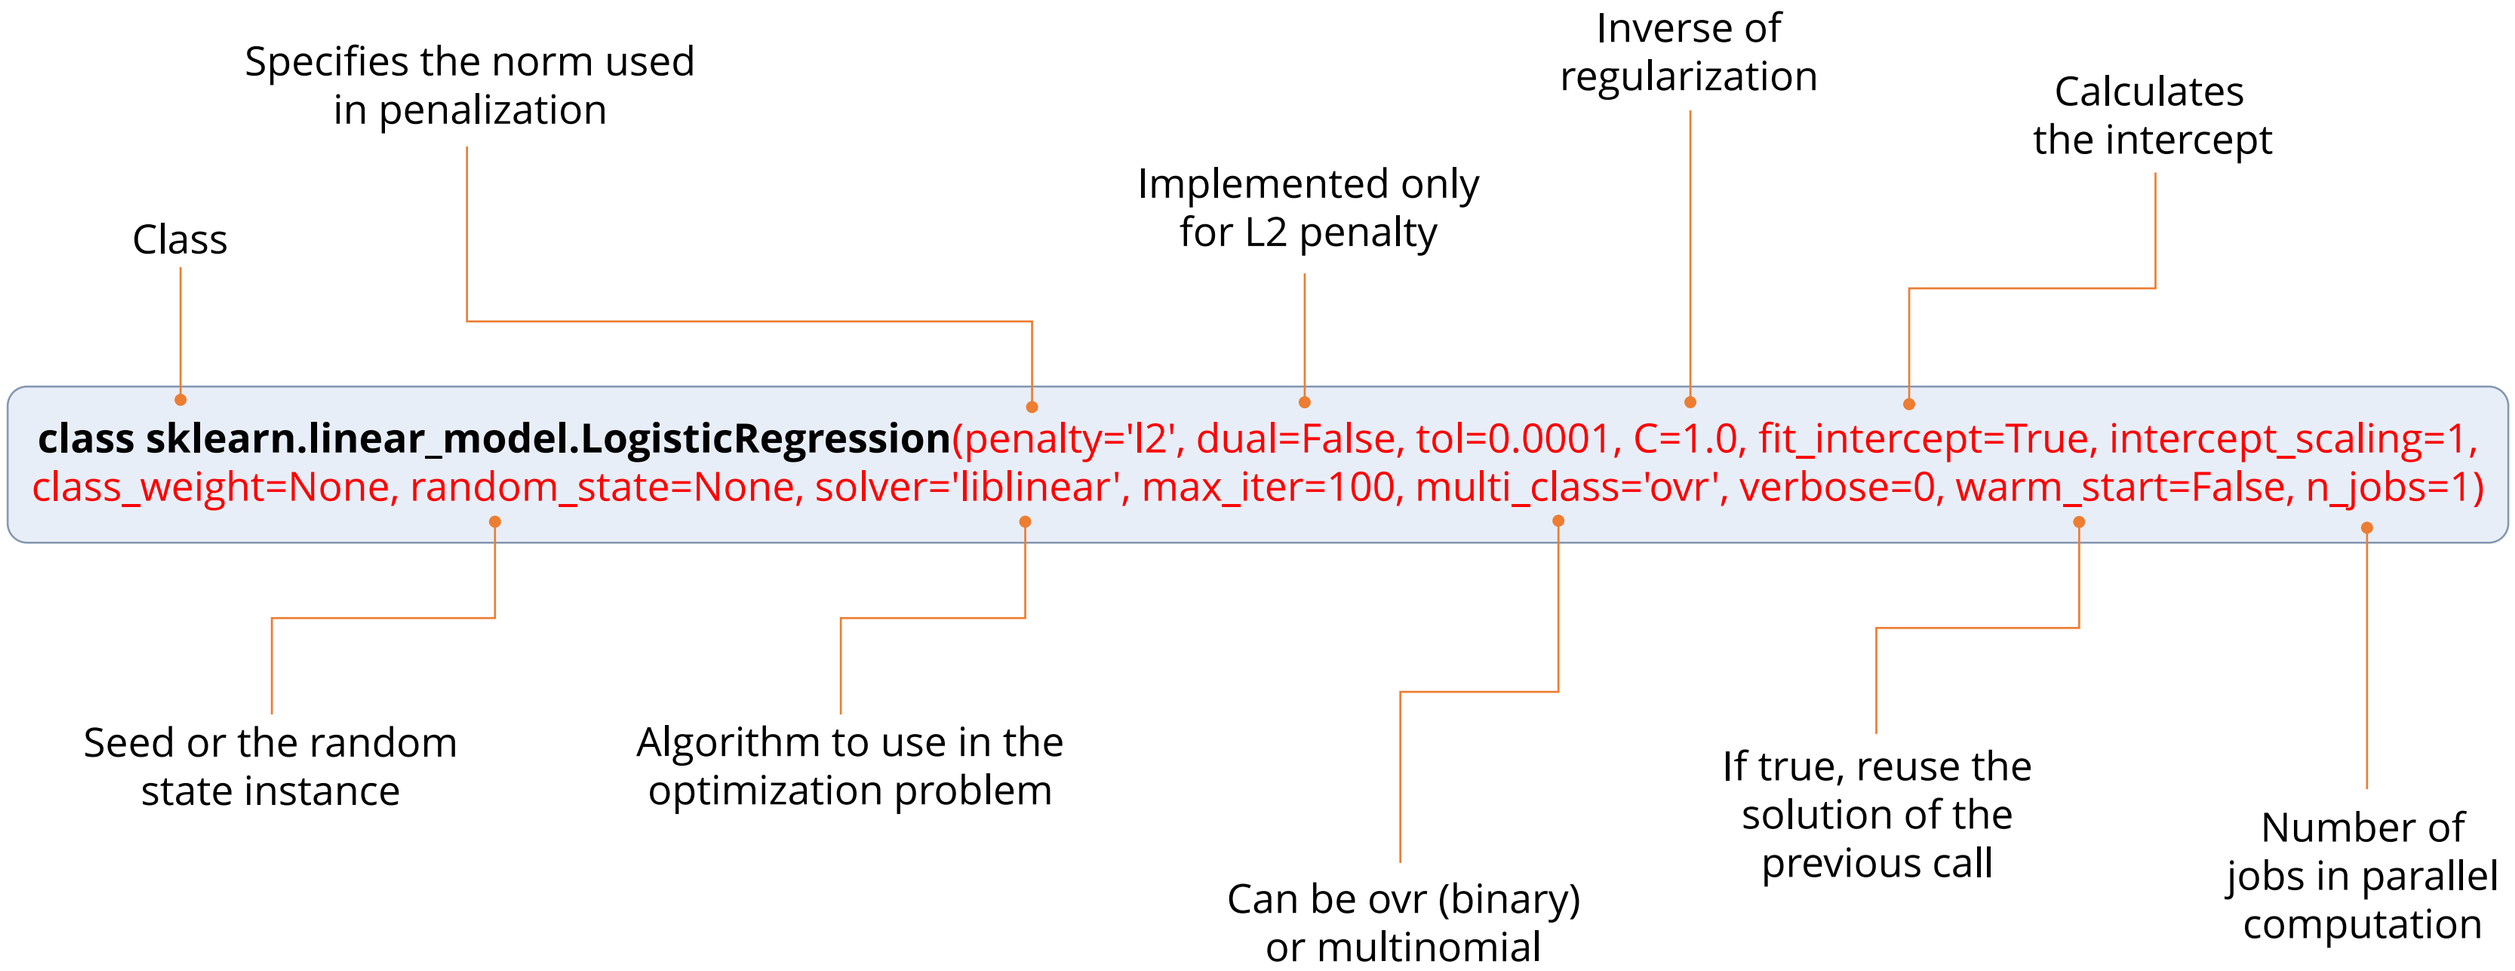

Supervised Learning Models: K Nearest Neighbors (K-NN)

K-nearest neighbors, or K-NN, is one of the simplest machine learning algorithms used for both classification and regression problem types.

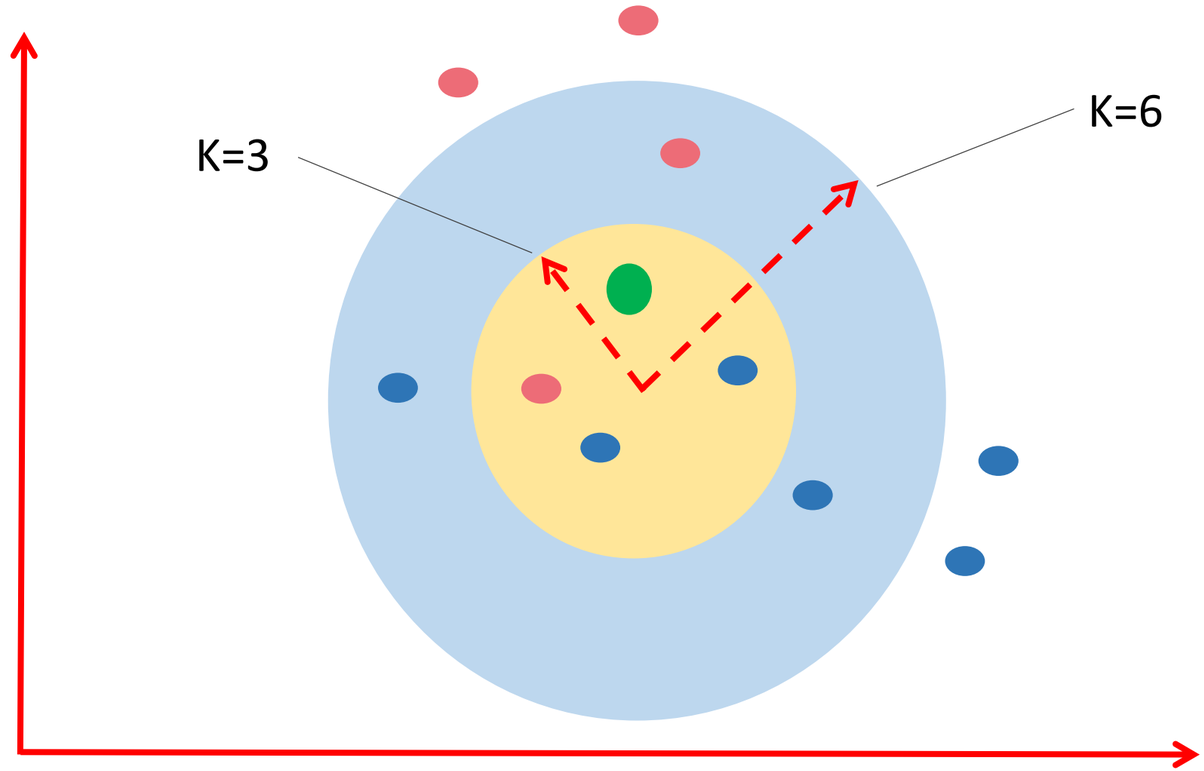

If you are using this method for binary classification, choose an odd number for k to avoid the case of a "tied" distance between two classes.

Demo 03—K-NN and Logistic Regression Models

Demonstrate the use of K-NN and logistic regression models

DATA SCIENCE

Unsupervised Learning Models: Clustering

A cluster is a group of similar data points.

It is used:

- To extract the structure of the data
- To identify groups in the data

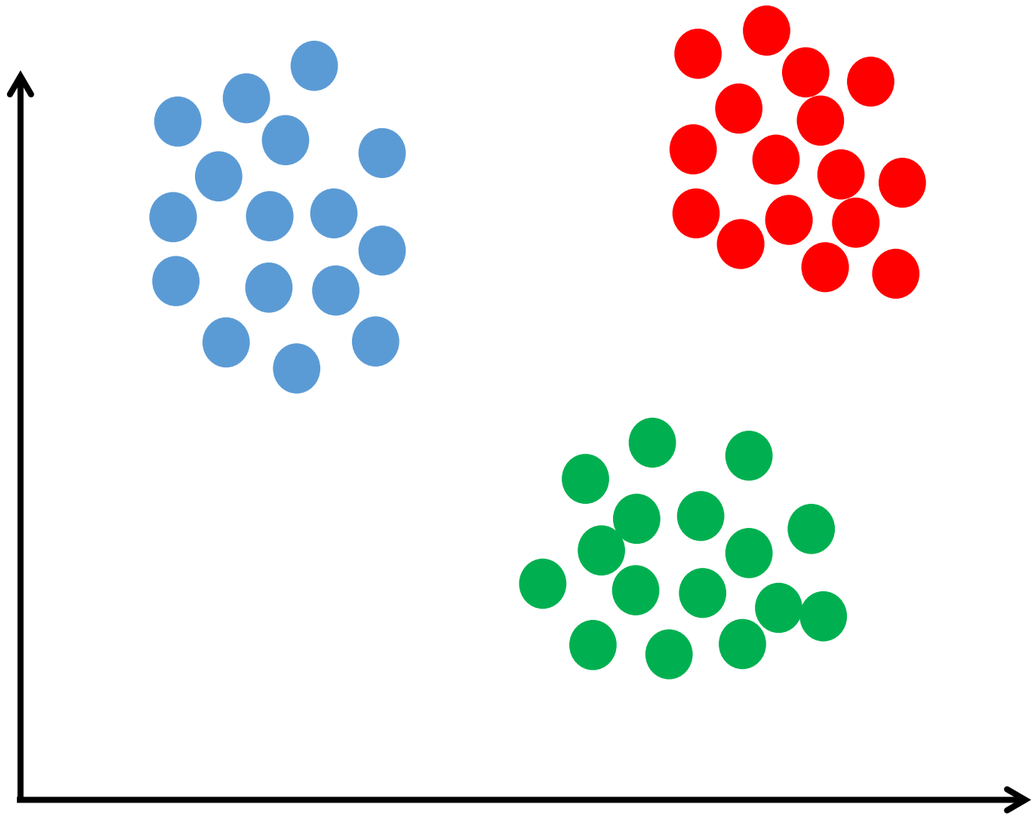

Greater similarity between data points results in better clustering.

Unsupervised Learning Models: K-means Clustering

K-means finds the best centroids by alternatively assigning random centroids to a dataset and selecting mean data points from the resulting clusters to form new centroids. It continues this process iteratively until the model is optimized.

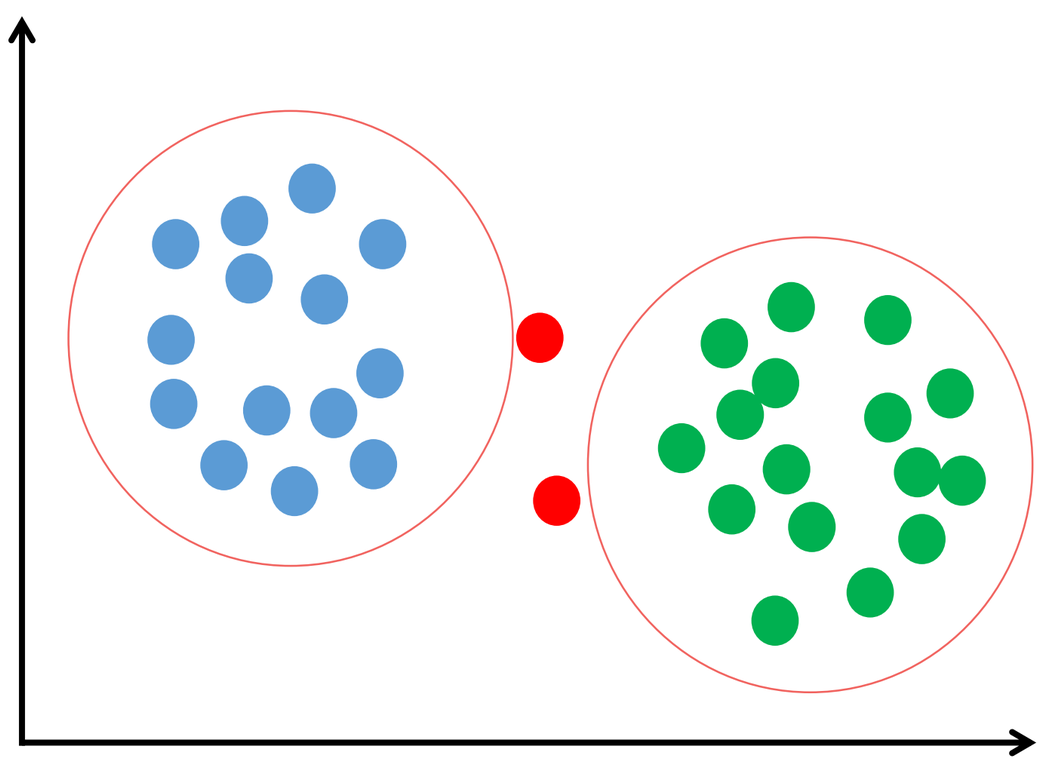

Assign

Find the number of clusters and assign mean

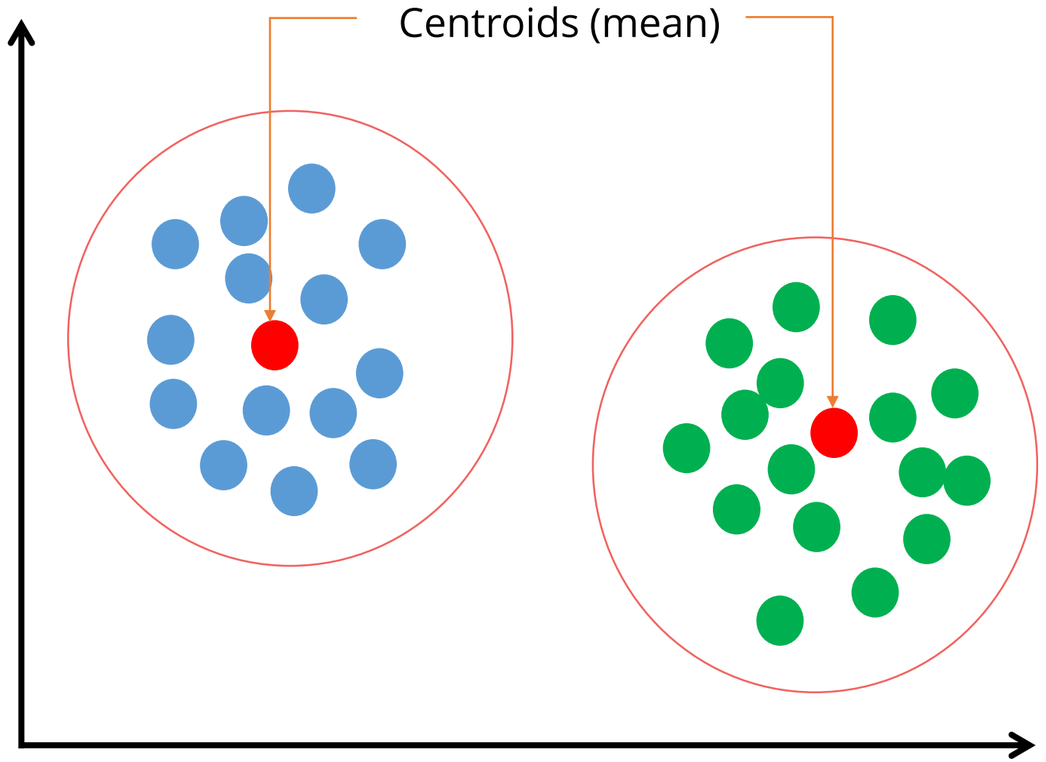

Optimize

Iterate and optimize the mean for each cluster for its respective data points

Unsupervised Learning Models: K-means Clustering (contd.)

Let us see how the k-means algorithm works in Scikit-Learn.

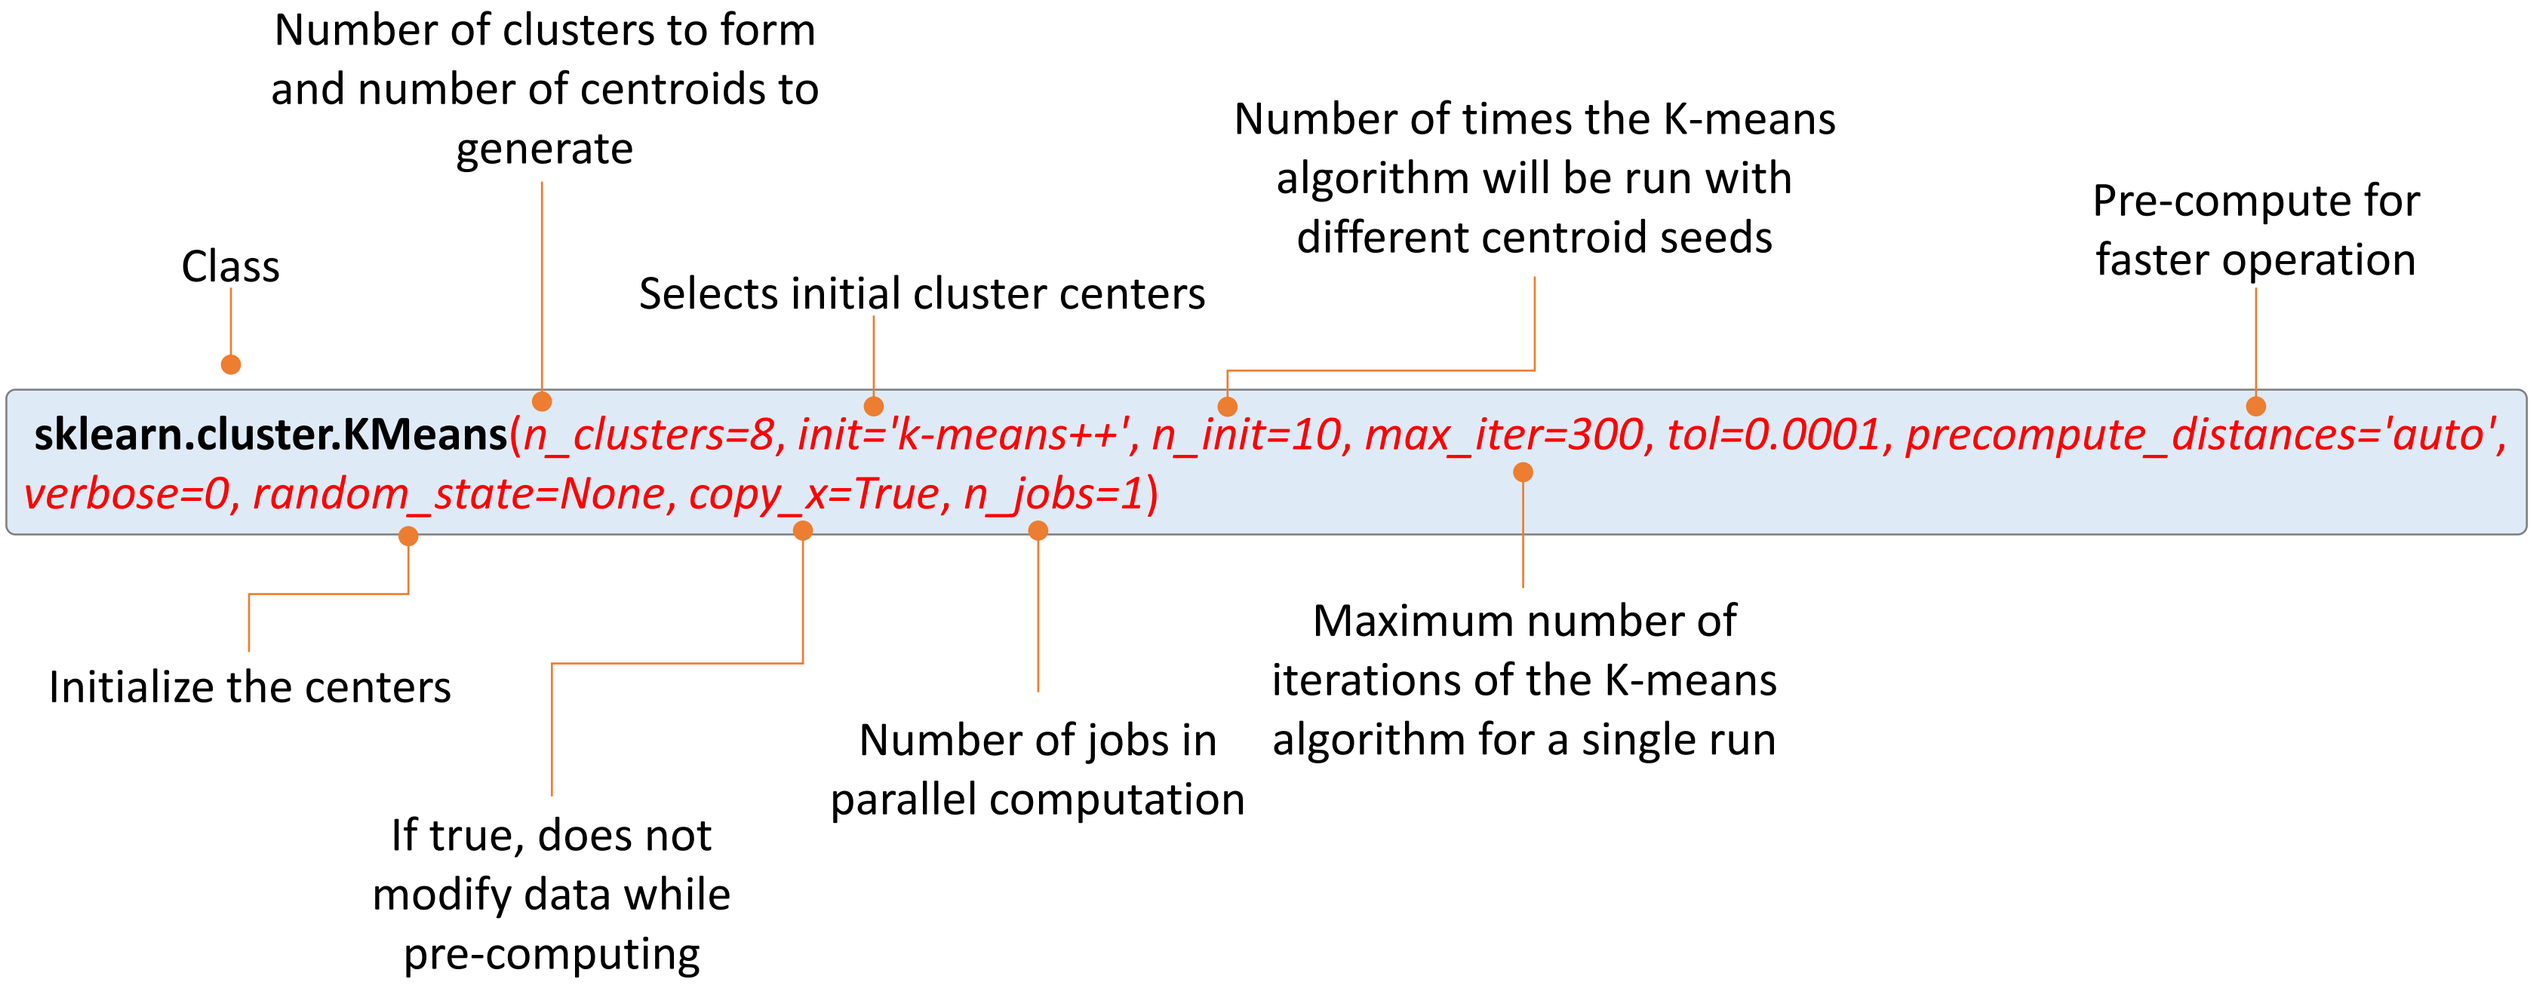

Demo 04—K-means Clustering

Demonstrate how to use k-means clustering to classify data points

DATA SCIENCE

Unsupervised Learning Models: Dimensionality Reduction

It reduces a high-dimensional dataset into a dataset with fewer dimensions. This makes it easier and faster for the algorithm to analyze the data.

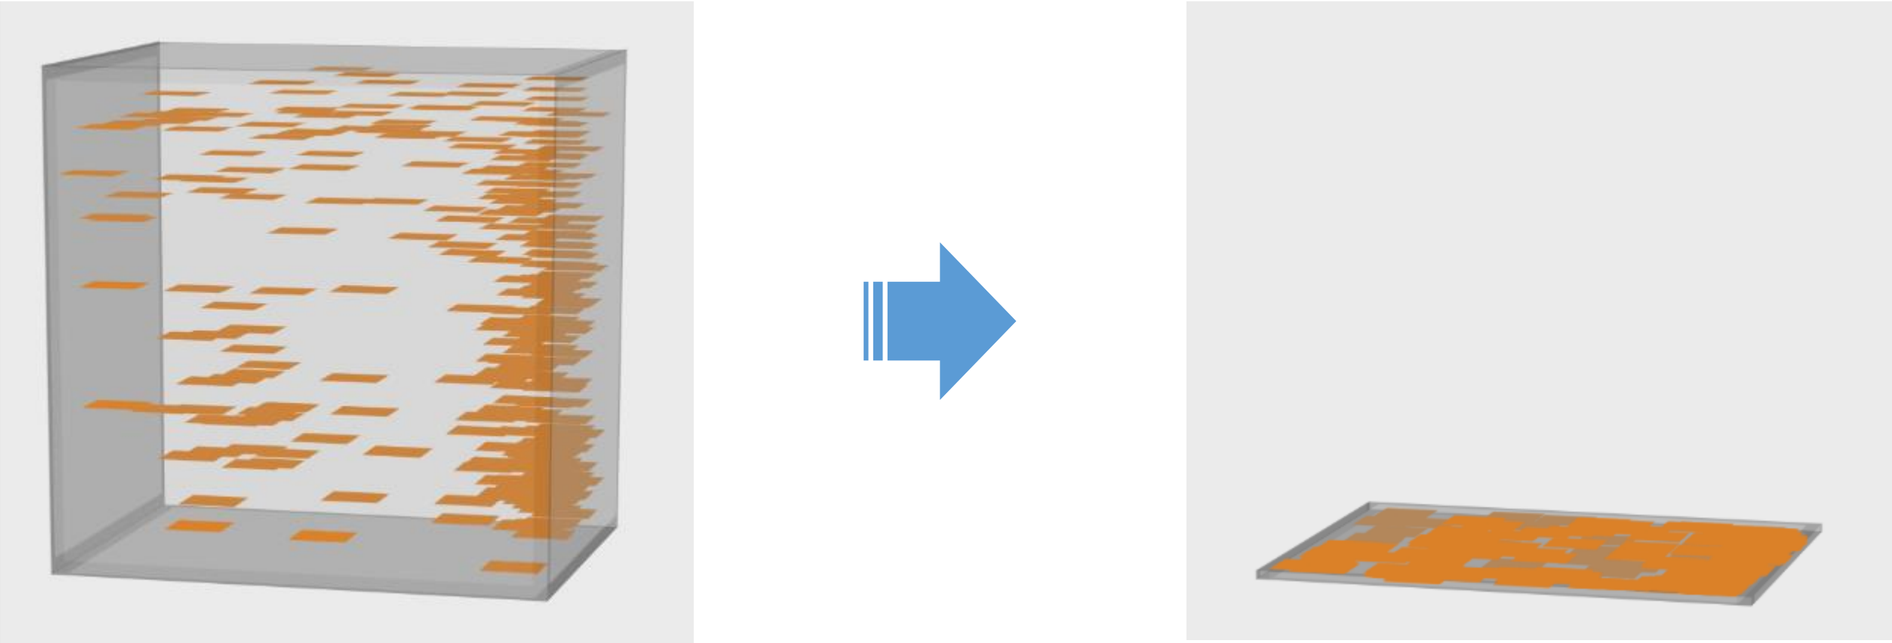

Unsupervised Learning Models: Dimensionality Reduction (contd.)

These are some techniques used for dimensionality reduction:

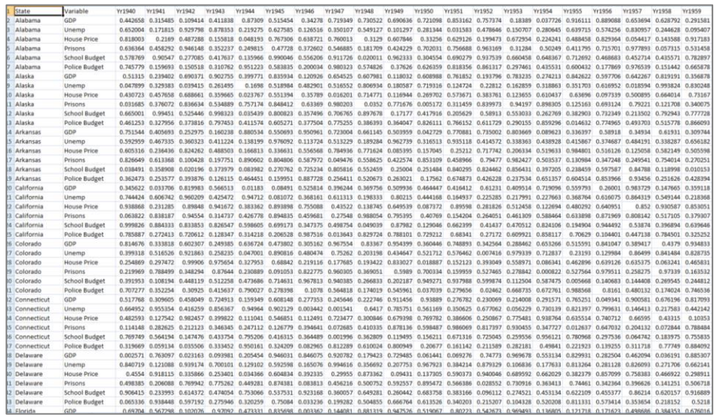

State	Variable	Yr1940	Yr1941	Yr1942	Yr1943	Yr1944	Yr1945	Yr1946	Yr1947	Yr1948	Yr1949	Yr1950	Yr1951	Yr1952	Yr1953	Yr1954	Yr1955	Yr1956	Yr1957	Yr1958	Yr1959
Alabama	GDP	0.442658	0.315485	0.109414	0.411838	0.87309	0.515454	0.34278	0.719349	0.730522	0.690636	0.721098	0.853162	0.757374	0.18389	0.037726	0.916111	0.889088	0.653694	0.628792	0.291581
Alabama	Unemp	0.652004	0.171815	0.929798	0.878353	0.219275	0.627585	0.126516	0.350107	0.549127	0.101297	0.281344	0.031583	0.478646	0.150707	0.280945	0.659715	0.336256	0.830957	0.244628	0.095407
Alabama	House Price	0.818003	0.2189	0.487288	0.155818	0.048193	0.767306	0.638721	0.760013	0.3129	0.607846	0.33256	0.629126	0.199473	0.672954	0.224241	0.488458	0.829364	0.054417	0.143588	0.917183
Alabama	Prisons	0.636364	0.458292	0.946146	0.352237	0.249815	0.47728	0.372602	0.546885	0.181709	0.424229	0.702031	0.756688	0.963169	0.31284	0.50249	0.411795	0.715701	0.977893	0.057313	0.531458
Alabama	School Budget	0.578769	0.90547	0.277085	0.417617	0.115966	0.990046	0.556206	0.911726	0.020011	0.962333	0.304554	0.690279	0.937539	0.660458	0.648367	0.712692	0.468683	0.452714	0.435571	0.782897
Alabama	Police Budget	0.745779	0.159693	0.150518	0.310762	0.951223	0.583835	0.200034	0.980323	0.574826	0.37626	0.626359	0.818356	0.861317	0.297461	0.435531	0.600432	0.117869	0.976539	0.151442	0.665878
Alaska	GDP	0.51315	0.239402	0.690371	0.902755	0.399771	0.835934	0.120926	0.654525	0.607981	0.118032	0.608988	0.761852	0.193796	0.783235	0.274213	0.842622	0.597706	0.642267	0.819191	0.356878
Alaska	Unemp	0.067899	0.325383	0.039415	0.265495	0.1698	0.518984	0.482901	0.518552	0.806934	0.180587	0.719516	0.124724	0.22812	0.162859	0.318863	0.351701	0.618952	0.018594	0.993824	0.830248
Alaska	House Price	0.430723	0.457658	0.668661	0.359665	0.023767	0.551394	0.35789	0.016201	0.714771	0.116944	0.269702	0.573671	0.383761	0.123655	0.610437	0.63696	0.097339	0.500895	0.664014	0.73167
Alaska	Prisons	0.031685	0.376072	0.836634	0.534889	0.757174	0.848412	0.63369	0.980203	0.0352	0.771676	0.005172	0.311459	0.839973	0.94197	0.898305	0.125163	0.691124	0.79221	0.121706	0.340075
Alaska	School Budget	0.665001	0.99451	0.525446	0.998123	0.035439	0.800823	0.357496	0.706765	0.897678	0.17177	0.417916	0.205629	0.58913	0.553033	0.262769	0.382903	0.732349	0.213002	0.792943	0.777728
Alaska	Police Budget	0.461253	0.327956	0.373816	0.797453	0.411574	0.605271	0.377504	0.775255	0.386393	0.364047	0.826111	0.766152	0.611729	0.290155	0.859296	0.014632	0.774965	0.493703	0.515778	0.866093
Arkansas	GDP	0.751544	0.405693	0.252975	0.160238	0.880534	0.550693	0.950961	0.723004	0.661145	0.503959	0.042729	0.770881	0.755002	0.803669	0.089623	0.336397	0.58918	0.34934	0.61931	0.309744
Arkansas	Unemp	0.592959	0.467155	0.366523	0.411224	0.138199	0.976092	0.113724	0.513229	0.185284	0.962739	0.316513	0.955118	0.414572	0.338363	0.438928	0.415867	0.378687	0.484191	0.338287	0.656182
Arkansas	House Price	0.605316	0.236436	0.824262	0.488503	0.166813	0.336631	0.556568	0.784936	0.771624	0.085395	0.157045	0.25212	0.717742	0.206334	0.519633	0.984801	0.516126	0.125058	0.582149	0.505598
Arkansas	Prisons	0.826649	0.611368	0.100428	0.197751	0.890602	0.804806	0.587972	0.049476	0.558625	0.422574	0.851109	0.458966	0.79477	0.983427	0.503537	0.130984	0.347248	0.249541	0.754014	0.270251
Arkansas	School Budget	0.038491	0.358908	0.020196	0.173979	0.083982	0.270762	0.725234	0.805816	0.552459	0.25004	0.251484	0.840295	0.824462	0.856431	0.397205	0.238459	0.597587	0.84788	0.118998	0.010153
Arkansas	Police Budget	0.362473	0.253577	0.393876	0.126115	0.464451	0.159951	0.887728	0.256411	0.520673	0.263021	0.17562	0.674873	0.426228	0.237534	0.651357	0.604514	0.853966	0.93456	0.251626	0.428394
California	GDP	0.345622	0.031706	0.819983	0.566513	0.01183	0.08491	0.525814	0.196244	0.369756	0.509936	0.464447	0.416412	0.81231	0.409514	0.719096	0.559793	0.26001	0.983729	0.147665	0.359118
California	Unemp	0.744424	0.606742	0.960209	0.425472	0.94712	0.081072	0.368161	0.611313	0.198333	0.80215	0.444168	0.1649317	0.225285	0.217991	0.227663	0.368764	0.616075	0.884419	0.549144	0.218368
California	House Price	0.938868	0.231285	0.89848	0.941672	0.383362	0.893898	0.755088	0.43522	0.138745	0.649539	0.087372	0.89598	0.281826	0.512458	0.122694	0.480292	0.640951	0.852	0.930587	0.850051
California	Prisons	0.063822	0.838187	0.94554	0.114737	0.426778	0.894835	0.459681	0.27548	0.988054	0.795395	0.40769	0.154204	0.264051	0.461309	0.588464	0.633898	0.871969	0.808142	0.517105	0.379307
California	School Budget	0.999826	0.884333	0.833853	0.826547	0.598605	0.699173	0.347175	0.498754	0.049039	0.87982	0.129046	0.662399	0.414137	0.470512	0.824106	0.194904	0.944492	0.53878	0.196894	0.639646
California	Police Budget	0.785887	0.272413	0.720612	0.128347	0.314218	0.206528	0.987516	0.013643	0.829724	0.788101	0.729212	0.68341	0.27172	0.609921	0.858117	0.70629	0.104401	0.447138	0.784501	0.325252
Colorado	GDP	0.814676	0.333818	0.602307	0.249385	0.636724	0.473802	0.305182	0.967554	0.83367	0.954399	0.360446	0.748893	0.342564	0.288462	0.653266	0.515591	0.841047	0.389417	0.4379	0.934833
Colorado	Unemp	0.199318	0.516526	0.921863	0.2582335	0.047001	0.890816	0.480474	0.75262	0.203198	0.444647	0.521712	0.570462	0.007416	0.979139	0.712837	0.23193	0.129984	0.86499	0.841484	0.828735
Colorado	House Price	0.254869	0.297472	0.99906	0.975654	0.327953	0.68842	0.219116	0.177685	0.193422	0.833027	0.018887	0.152123	0.393049	0.558971	0.086341	0.462896	0.639126	0.635575	0.063241	0.465831
Colorado	Prisons	0.219969	0.788499	0.348294	0.87644	0.230889	0.091053	0.822775	0.960305	0.369051	0.5989	0.700334	0.159959	0.527465	0.278842	0.000822	0.527564	0.979511	0.258275	0.97339	0.163532
Colorado	School Budget	0.391953	0.108194	0.448119	0.312258	0.473686	0.714631	0.967813	0.940385	0.266833	0.202187	0.945271	0.917988	0.599874	0.112504	0.587475	0.005668	0.140683	0.144408	0.2695545	0.244812
Colorado	Police Budget	0.707277	0.352254	0.30925	0.415637	0.790027	0.278398	0.1078	0.564818	0.174019	0.545961	0.037039	0.279656	0.02462	0.668735	0.672761	0.988568	0.8161	0.480132	0.174024	0.746536
Connecticut	GDP	0.517768	0.309605	0.458049	0.724913	0.159349	0.608148	0.277353	0.245646	0.222746	0.911456	0.93889	0.276782	0.230069	0.214008	0.291571	0.765251	0.049341	0.900581	0.676196	0.817093
Connecticut	Unemp	0.664952	0.955354	0.416259	0.8563867	0.94964	0.902129	0.003442	0.001541	0.6417	0.785751	0.561689	0.350625	0.677062	0.056229	0.730139	0.821397	0.799631	0.146413	0.217583	0.442142
Connecticut	House Price	0.482593	0.127542	0.982457	0.399822	0.111041	0.546851	0.112491	0.723477	0.300846	0.679396	0.769782	0.386606	0.250867	0.775481	0.938764	0.635514	0.740712	0.66595	0.43315	0.10353
Connecticut	Prisons	0.114148	0.282625	0.212123	0.346345	0.247112	0.126779	0.394641	0.072685	0.410335	0.878136	0.598487	0.980069	0.817397	0.930455	0.347727	0.012637	0.647032	0.204132	0.072844	0.788484
Connecticut	School Budget	0.769749	0.564194	0.187476	0.413754	0.795206	0.416515	0.364489	0.002996	0.362809	0.119495	0.156211	0.671316	0.725045	0.259556	0.956121	0.780968	0.297536	0.064742	0.183975	0.755835
Connecticut	Police Budget	0.319669	0.059134	0.035506	0.333452	0.950161	0.324209	0.082965	0.812289	0.610024	0.800949	0.20677	0.161142	0.211589	0.282181	0.49841	0.221923	0.139255	0.311718	0.77749	0.884092
Delaware	GDP	0.002571	0.763097	0.023163	0.093981	0.205454	0.946031	0.846675	0.920782	0.179423	0.729485	0.061441	0.089276	0.74773	0.969678	0.553134	0.829931	0.282504	0.462094	0.036191	0.885307
Delaware	Unemp	0.840719	0.121088	0.939174	0.700101	0.129102	0.592598	0.165076	0.994616	0.356692	0.207753	0.967923	0.384214	0.879129	0.106836	0.177633	0.811264	0.281128	0.826091	0.271706	0.662141
Delaware	House Price	0.4554	0.918115	0.335866	0.253401	0.034366	0.604834	0.392335	0.29955	0.873362	0.09431	0.117305	0.590373	0.940046	0.689592	0.662029	0.382379	0.857099	0.758383	0.466922	0.298911
Delaware	Prisons	0.498385	0.206088	0.789942	0.775262	0.849281	0.874381	0.083813	0.456218	0.500752	0.592557	0.895472	0.566386	0.028552	0.730916	0.363413	0.74461	0.342364	0.396826	0.141251	0.506718
Delaware	School Budget	0.909415	0.233955	0.614372	0.474054	0.753066	0.537551	0.923168	0.360057	0.449281	0.260442	0.683758	0.383166	0.096112	0.274521	0.453134	0.622109	0.455377	0.86214	0.620157	0.916889
Delaware	Police Budget	0.065336	0.938448	0.597192	0.275946	0.320259	0.75084	0.033236	0.199282	0.504855	0.666764	0.613526	0.340203	0.215207	0.104328	0.520208	0.811331	0.573414	0.353654	0.218152	0.5218
Florida	GDP	0.69204	0.567298	0.103026	0.92092	0.473331	0.835698	0.003367	0.144081	0.881319	0.947526	0.519067	0.80721	0.542673	0.969493	0.136805	0.121718	0.172623	0.498686	0.384351	0.629016

Large dataset
(a few thousand columns and rows)

- Drop data columns with missing values
- Drop data columns with low variance
- Drop data columns with high correlations
- Apply statistical functions - PCA

Unsupervised Learning Models: Principal Component Analysis (PCA)

It is a linear dimensionality reduction method which uses singular value decomposition of the data and keeps only the most significant singular vectors to project the data to a lower dimensional space.

- It is primarily used to compress or reduce the data.
- PCA tries to capture the variance, which helps it pick up interesting features.
- PCA is used to reduce dimensionality in the dataset and to build our feature vector.
- Here, the principal axes in the feature space represents the direction of maximum variance in the data.

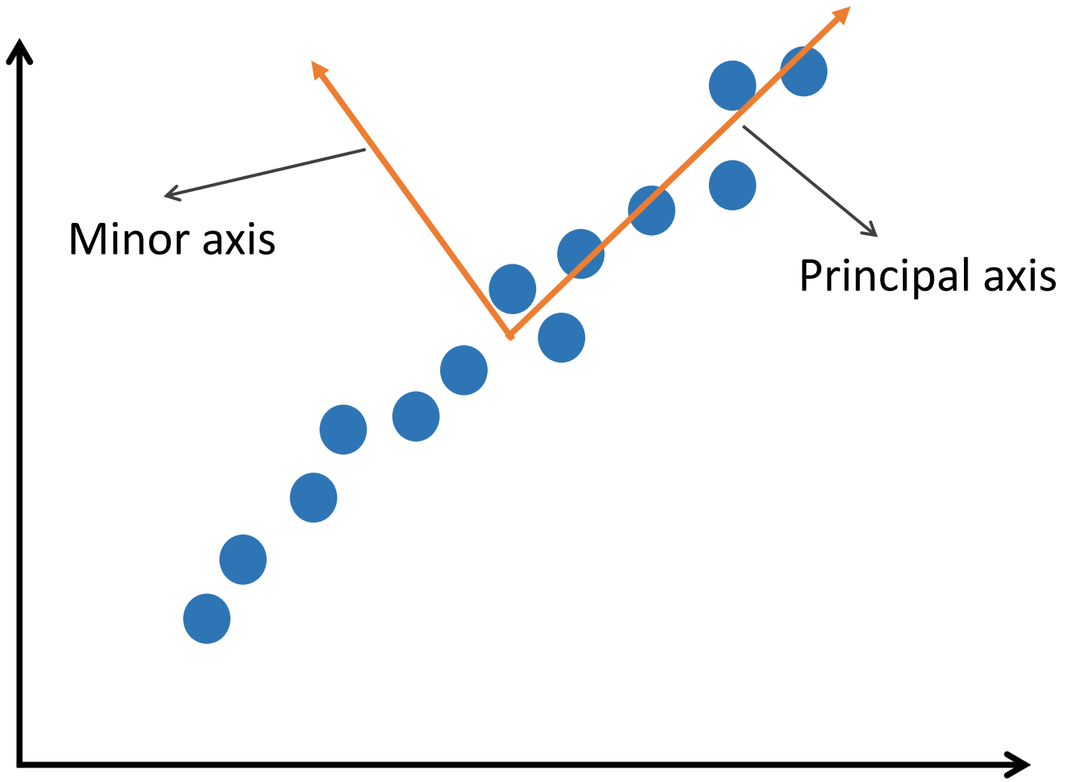

This method is used to capture variance.

Unsupervised Learning Models: Principal Component Analysis (PCA)

Let us look at how the PCA algorithm works in Scikit-Learn.

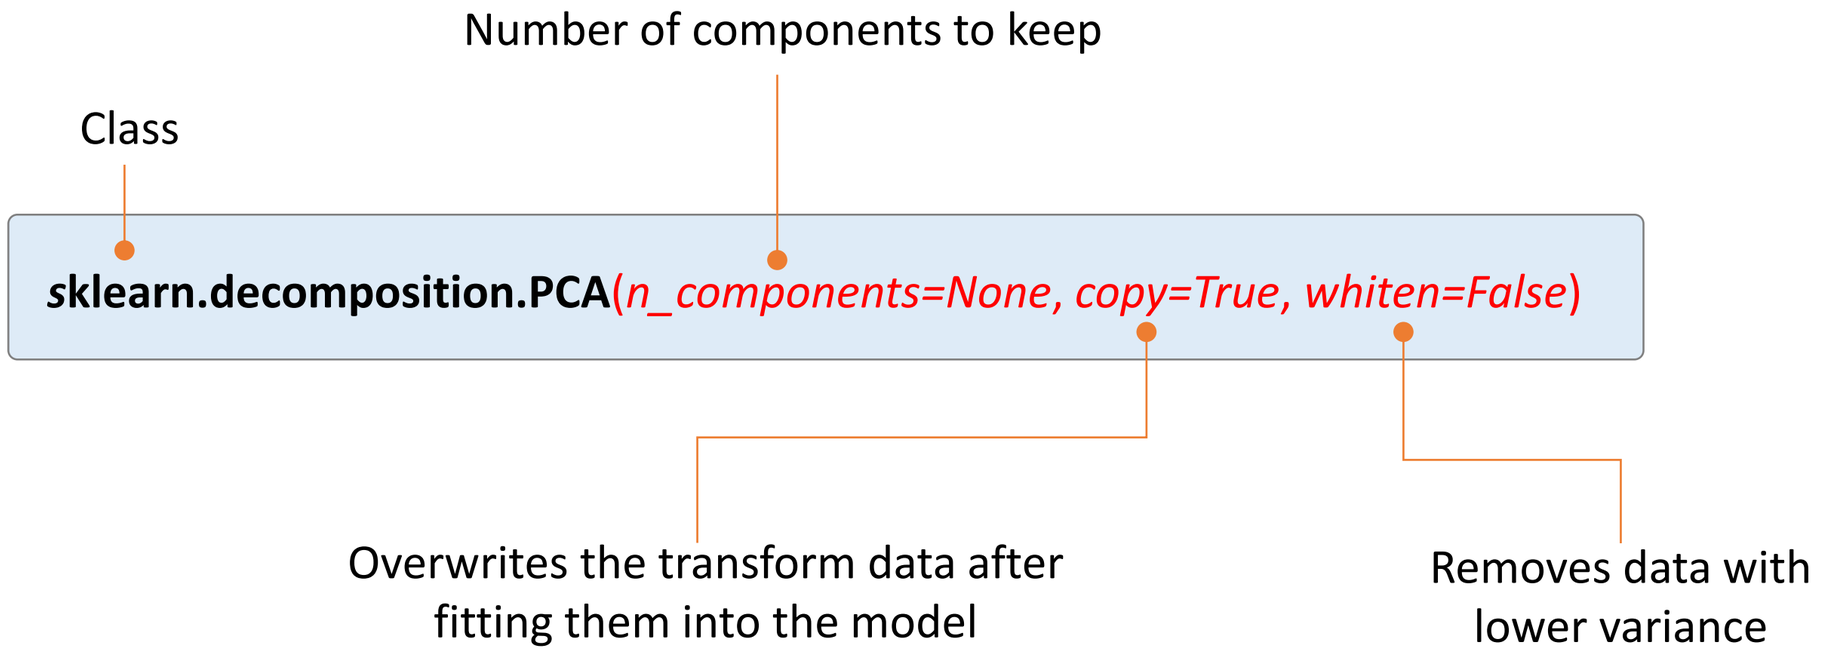

Demo 05—Principal Component Analysis (PCA)

Demonstrate how to use the PCA model to reduce the dimensions of a dataset

DATA SCIENCE

Pipeline

Pipeline is mainly used to combine multiple models or estimators. Its characteristics are as follows:

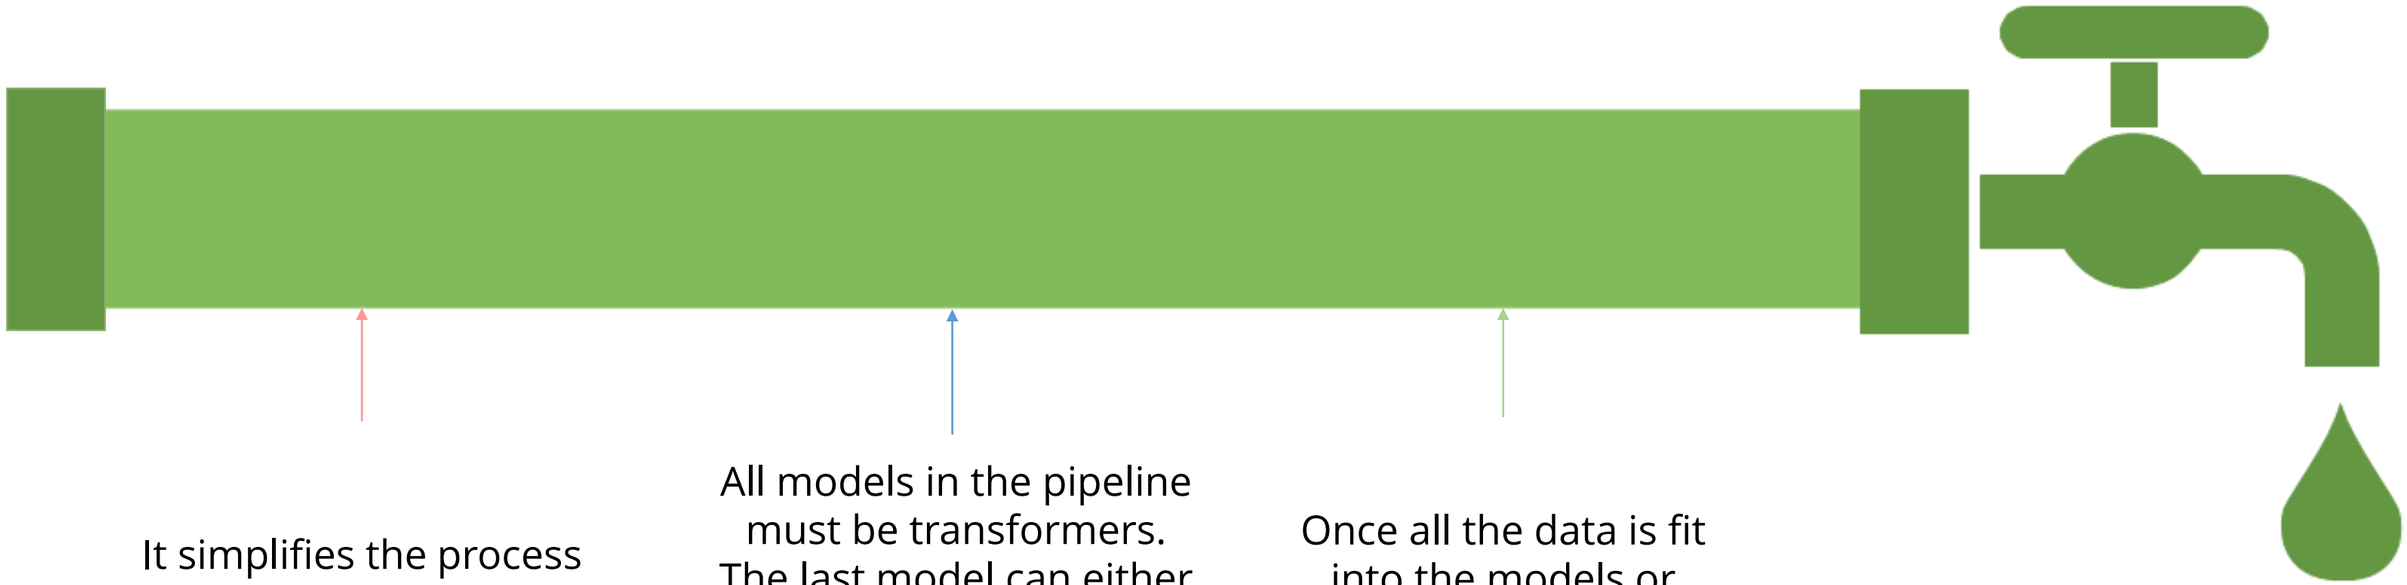

It simplifies the process where more than one model is required or used.

All models in the pipeline must be transformers. The last model can either be a transformer or a classifier, regressor, or other such objects.

Once all the data is fit into the models or estimators, the predict method can be called.

Estimators are known as ‘model instance’.

Demo 06—Pipeline

Demonstrate how to build a pipeline

DATA SCIENCE

Model Persistence

Save model for the future use. No need to retrain your model every time when you need them.

It is possible to save a model by using Python's Pickle method.

Scikit-learn has a special replacement for pickle called joblib.

You can use joblib.dump and joblib.load methods.

These are more efficient for Big Data.

Demo 06—Model Persistence

Demonstrate how to persist a model for future use

DATA SCIENCE

Model Evaluation: Metric Functions

You can use the “Metrics” function to evaluate the accuracy of your model’s predictions.

Classification	`metrics.accuracy_score` `metrics.average_precision_score`
Clustering	`metrics.adjusted_rand_score`

Regression	`metrics.mean_absolute_error`
	`metrics.mean_squared_error`
	`metrics.median_absolute_error`

Knowledge Check

KNOWLEDGE CHECK

What is the best way to train a model?

a. Use the entire dataset as a training and testing set

b. Split the known dataset into separate training and testing sets

c. Ask the source to provide continuous data

d. Ask the source to provide categorical data

KNOWLEDGE CHECK

What is the best way to train a model?

a. Use the entire dataset as a training and testing set both

b. Split the known dataset into separate training and testing sets

c. Ask the source to provide continuous data

d. Ask the source to provide categorical data

The correct answer is **b**.

Explanation: The best way to train a model is to split the known dataset into training and testing sets. The testing set varies from 20% to 40%.

Assignment

Assignment 01

Problem | Instructions

The given dataset contains ad budgets for different media channels and the corresponding ad sales of XYZ firm. Evaluate the dataset to:

- Find the features or media channels used by the firm
- Find the sales figures for each channel
- Create a model to predict the sales outcome
- Split as training and testing datasets for the model
- Calculate the Mean Square Error (MSE)

Assignment 01

Problem | Instructions

Instructions on performing the assignment:

- Download the “Advertising Budget and Sales.csv” file from the “Resource” tab. You can load the saved file to the Jupyter notebook that you would be using to complete the assignment.

Common instructions:

- If you are new to Python, download the “Anaconda Installation Instructions” document from the “Resources” tab to view the steps for installing Anaconda and the Jupyter notebook.
- Download the “Assignment 01” notebook and upload it on the Jupyter notebook to access it.
- Follow the provided cues to complete the assignment.

Assignment

Assignment 02

Problem | Instructions

The given dataset lists the glucose level readings of several pregnant women taken either during a survey examination or routine medical care. It specifies if the 2 hour post-load plasma glucose was at least 200 mg/dl. Analyze the dataset to:

1. Find the features of the dataset,
2. Find the response label of the dataset,
3. Create a model to predict the diabetes outcome,
4. Use training and testing datasets to train the model, and
5. Check the accuracy of the model.

Assignment 02

Problem | Instructions

Instructions on performing the assignment:

- Download the “pima-indians-diabetes.DATA” and “pima-indians-diabetes.NAMES” files from the “Resources” tab. Load the .DATA file to the Jupyter notebook to work on it.
- Open the .NAMES file with a notepad application to view its text. Use this file to view the features of the dataset and add them manually in your code.

Common instructions:

- If you are new to Python, download the “Anaconda Installation Instructions” document from the “Resources” tab to view the steps for installing Anaconda and the Jupyter notebook.
- Download the “Assignment 01” notebook and upload it on the Jupyter notebook to access it.
- Follow the provided cues to complete the assignment.

Quiz

QUIZ 1

Which of the following is true with a greater value of SSR or SSE? *Select all that apply*.

a. The prediction will be more accurate, making it the best fit model.

b. The prediction will start becoming less accurate.

c. The outcome remains unaffected.

d. The model will not be the best fit for the attributes.

QUIZ 1

Which of the following is true with a greater value of SSR or SSE? Select all that apply.

a. The prediction will be more accurate, making it the best fit model.

b. The prediction will start becoming less accurate.

c. The outcome remains unaffected.

d. The model will not be the best fit for the attributes.

The correct answer is **b, d**.

Explanation: With higher SSR or SSE, the prediction will be less accurate and the model will not be the best fit for the attributes.

QUIZ 2

Class sklearn.linear_model.LogisticRegression, random_state _____.

a. indicates the seed of the pseudo random number generator used to shuffle data

b. defines the features state

c. represents the number of random iterations

d. specifies a random constant to be added to the decision function

QUIZ

2

Class sklearn.linear_model.LogisticRegression, random_state _____.

a. indicates the seed of the pseudo random number generator used to shuffle data

b. defines the features state

c. represents the number of random iterations

d. specifies a random constant to be added to the decision function

The correct answer is **a**.

Explanation: The class “sklearn.linear_model.LogisticRegression, random_state” indicates the seed of the pseudo random number generator used to shuffle data.

QUIZ 3

What are the requirements of the K-means algorithm? *Select all that apply*.

a. Number of clusters should be specified

b. More than one iteration should meet requisite criteria

c. Centroids should minimize inertia

d. Features should be labeled

QUIZ 3

What are the requirements of the K-means algorithm? *Select all that apply*.

a. Number of clusters should be specified

b. More than one iteration should meet requisite criteria

c. Centroids should minimize inertia

d. Features should be labeled

The correct answer is **a, b, c**.

Explanation: The K-means algorithm requires that the number of clusters be specified and that centroids that minimize inertia be selected. It requires several iterations to fine tune itself and meet the required criteria to become the best fit model.

QUIZ

4

In Class sklearn.decomposition.PCA, the transform(X) method , where X is multi-dimensional _____.

a. fits the model with X and applies the dimensionality reduction on X

b. transforms the data back to its original space

c. applies the dimensionality reduction on X

d. computes data co-variance with the generative model

QUIZ

3

In Class sklearn.decomposition.PCA, the transform(X) method , where X is multi-dimensional _____.

a. fits the model with X and applies the dimensionality reduction on X

b. transforms the data back to its original space

c. applies the dimensionality reduction on X

d. computes data co-variance with the generative model

The correct answer is **c.**

Explanation: In Class "sklearn.decomposition.PCA," the transform(X) method applies the dimensionality reduction on X.

Key Takeaways

- Scikit-learn has many built-in functions and algorithms which make it easy for Data Scientists to build machine learning models.
- Machine learning can easily and quickly extract information from large sources of data.
- Supervised and unsupervised machine learning models are two of the most widely used learning models.
- Supervised learning models are used to predict the outcome of a dataset.
- Unsupervised learning models are used to find the structure of a dataset.
- For continuous data, you can use regression algorithms if it is a supervised learning model. However, use dimensionality reduction for unsupervised learning.
- For categorical data, use classification algorithms if it is a supervised learning model and clustering if it is an unsupervised learning model.

This concludes “Machine Learning with Scikit-Learn.”

The next lesson is “Natural Language Processing (NLP) with Scikit-Learn.”